

Bias, Efficiency, and the Gauss Markov Theorem

Ben Gilbert

9/1/2021

Bias, Efficiency, and the Gauss Markov Theorem

- This is a program to illustrate bias and inefficiency when Gauss Markov assumptions fail
 - Bias: deviation of expected value of sample parameter estimate from “true” population parameter
 - Efficiency: the variance of the sample parameter estimate should be as small as possible
 - Consistency: distribution of sample parameter estimate should converge to population value as sample size grows
- Gauss-Markov: OLS (Ordinary Least Squares) is Best (lowest variance) Linear Unbiased Estimator (BLUE) IF:

1. true model linear in parameters and residuals:

$$y_t = \beta_0 + \beta_1 * x_{1t} + \beta_2 * x_{2t} + e_t$$

2. X variables (right hand side) are not constants or perfectly correlated with each other
3. Residuals “e” have constant variance
 - (not more noisy for some X’s than others, homoskedasticity vs. heteroskedasticity)
4. Residuals “e” are uncorrelated with each other
 - (no peer effects, no serial correlation)
5. All X variables are uncorrelated with the residual e
 - (observed X is not picking up some unobserved or uncontrolled factor)

- In each case we will run the linear regression

$$y_t = \beta_0 + \beta_1 * x_{1t} + \beta_2 * x_{2t} + e_t$$

on the data, but the “true” model or “data generating process” is different

Define directory and load packages

Load (and install if necessary) any packages that we want to use. If we want to define a working directory for this session, we can.

```
setwd("C:/Users/bgilbert_a/Dropbox/Econometrics/TimeSeriesCourse/Fall2021")
# install.packages("MASS")
require(MASS)
# install.packages("car")
require(car)
```

Violate Assumption 1: True model is not linear in parameters

- Set the seed for replicability.
- Generate the true model

```
set.seed(826)
# Covariance matrix of x1, x2, and e
Sig <- matrix(c(4,1,0,1,2,0,0,0,1),3,3)
Sig      # Notice e does not covary with x1 or x2 (assumption 5)

##      [,1] [,2] [,3]
## [1,]    4    1    0
## [2,]    1    2    0
## [3,]    0    0    1

      # also x1 and x2 can covary, but not perfectly (assumption 2)

# Mean of x1, x2, and e
moo <- c(10,3,0)

# generate data
Xe <- mvrnorm(n=1000,mu=moo,Sigma=Sig)

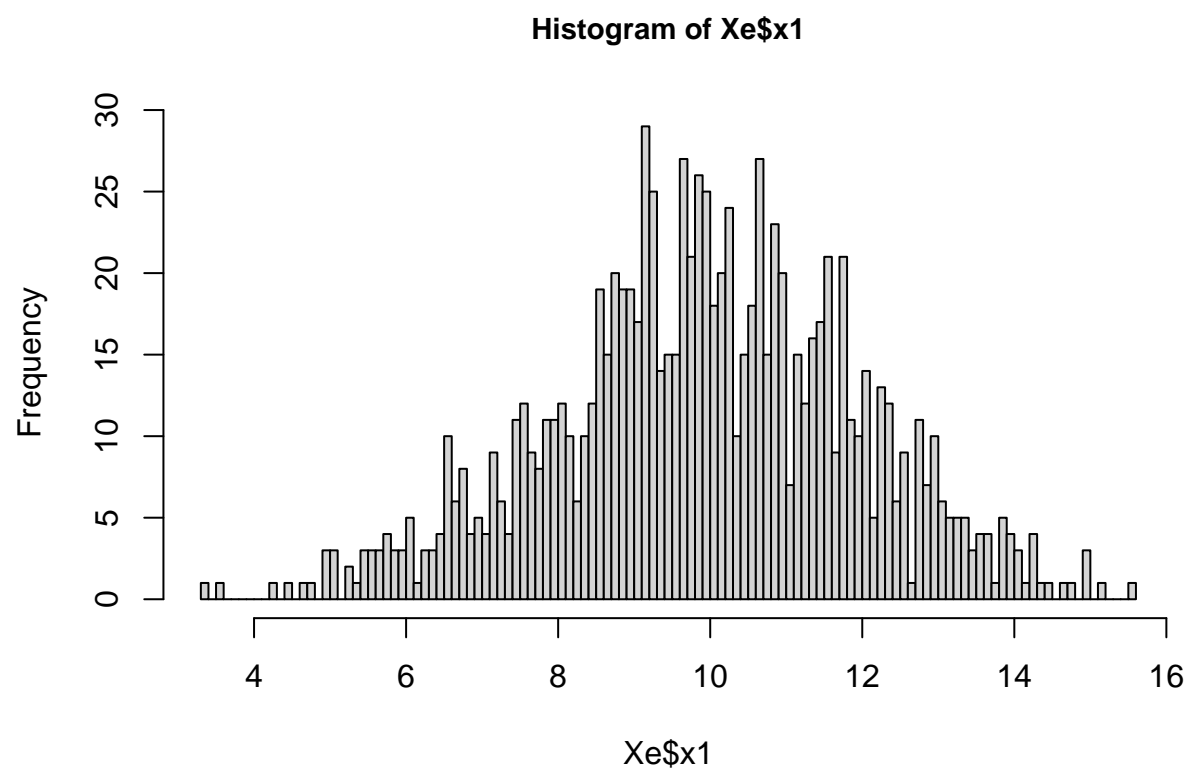
# give the variables names
colnames(Xe)<-c("x1", "x2", "e")

# store as a data frame
Xe <- as.data.frame(Xe)
head(Xe)

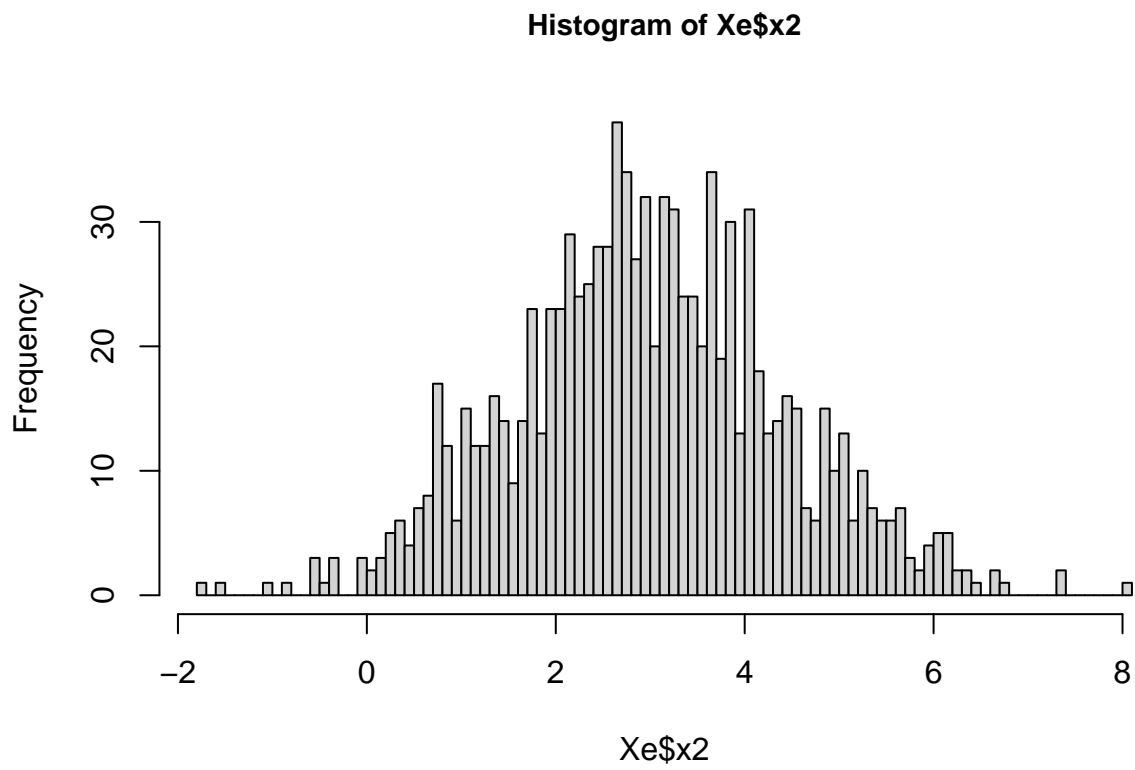
##      x1      x2      e
## 1  9.815559 2.8349325 1.0499626
## 2  8.443130 2.1208185 0.7884180
## 3  9.618013 3.3454521 -0.8777746
## 4  6.371419 1.7631024 0.6244005
## 5 10.091449 0.6709607 0.2481989
## 6  5.828882 1.1333412 0.4228365
```

- Investigate the data.
 - plot empirical distribution of each variable

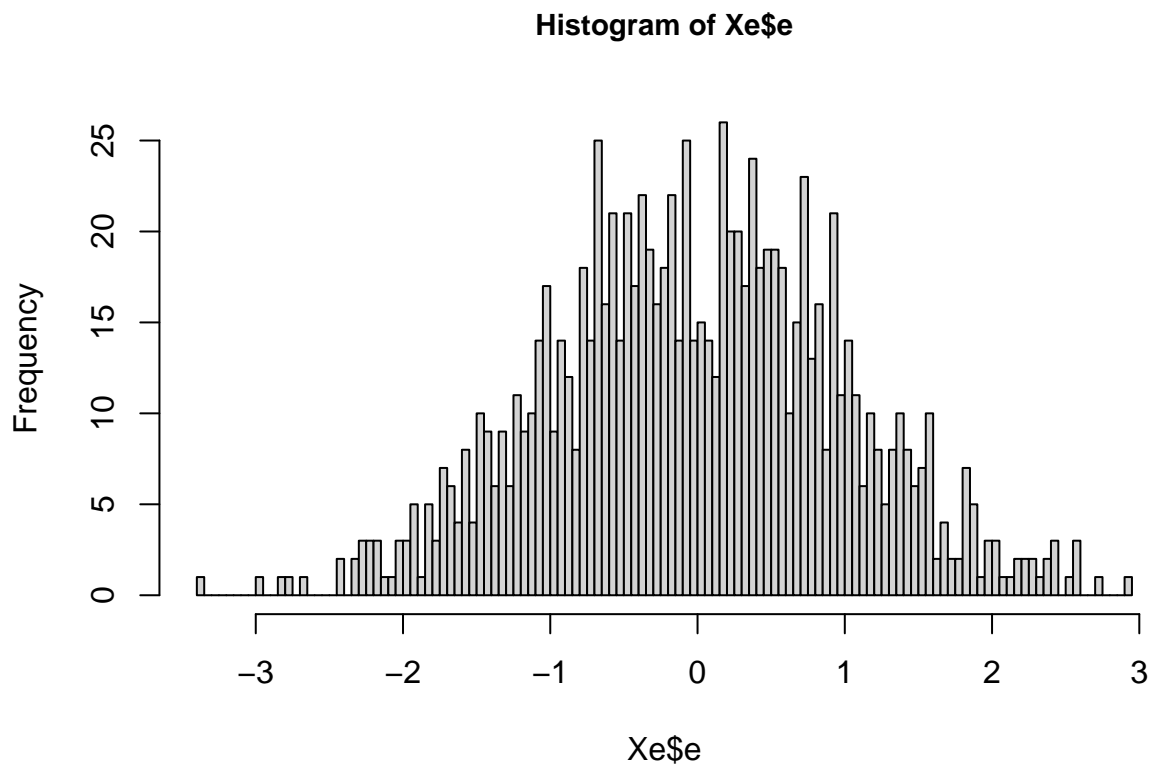
```
hist(Xe$x1, breaks = 100, cex.main = 0.9)
```



```
hist(Xe$x2, breaks = 100, cex.main = 0.9)
```



```
hist(Xe$e, breaks = 100, cex.main = 0.9)
```



Sample correlations and covariances (notice difference from "truth")

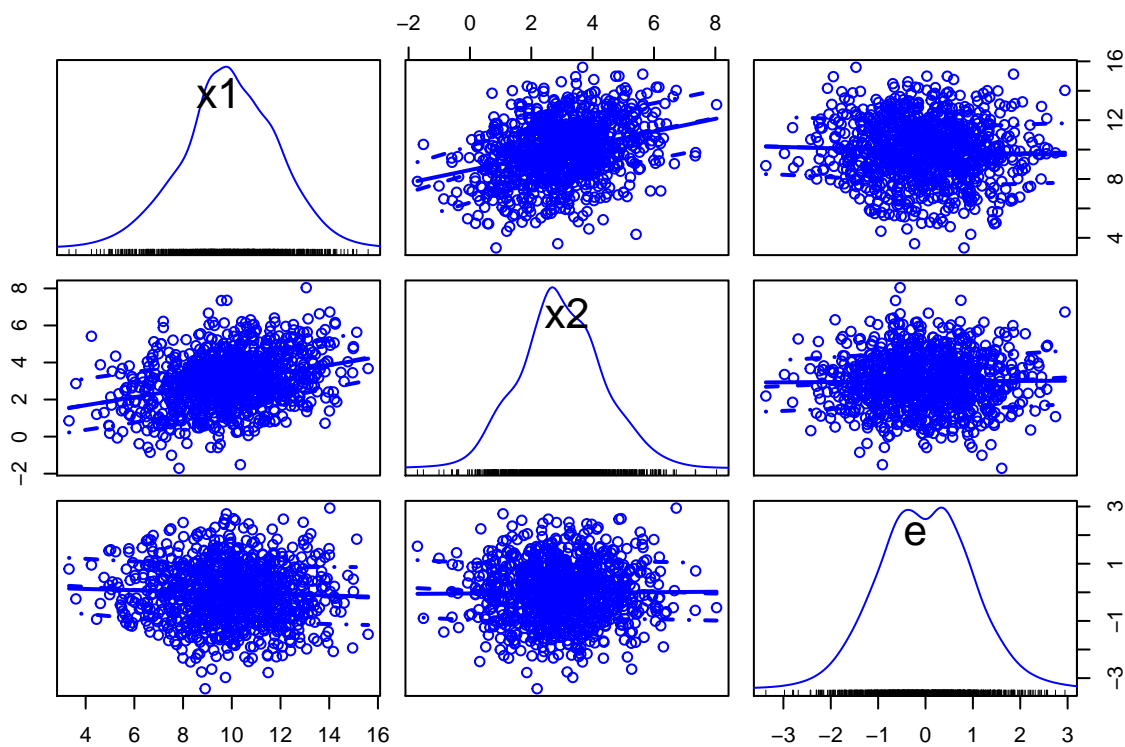
```
cov(Xe)
```

```
##           x1           x2           e
## x1  4.02081193  0.87289399 -0.09254668
## x2  0.87289399  2.00633796  0.01490771
## e   -0.09254668  0.01490771  1.00873015
```

```
cor(Xe)
```

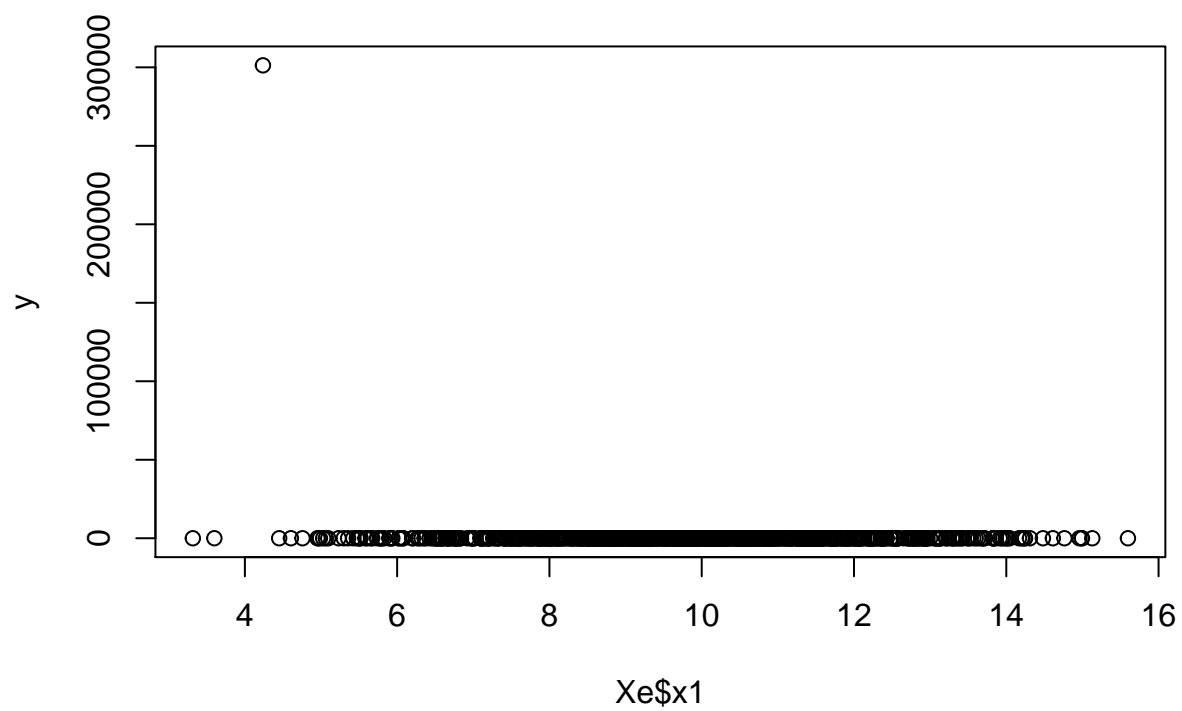
```
##           x1           x2           e
## x1  1.00000000  0.30732832 -0.04595327
## x2  0.30732832  1.00000000  0.01047904
## e   -0.04595327  0.01047904  1.00000000
```

```
scatterplotMatrix(Xe)
```

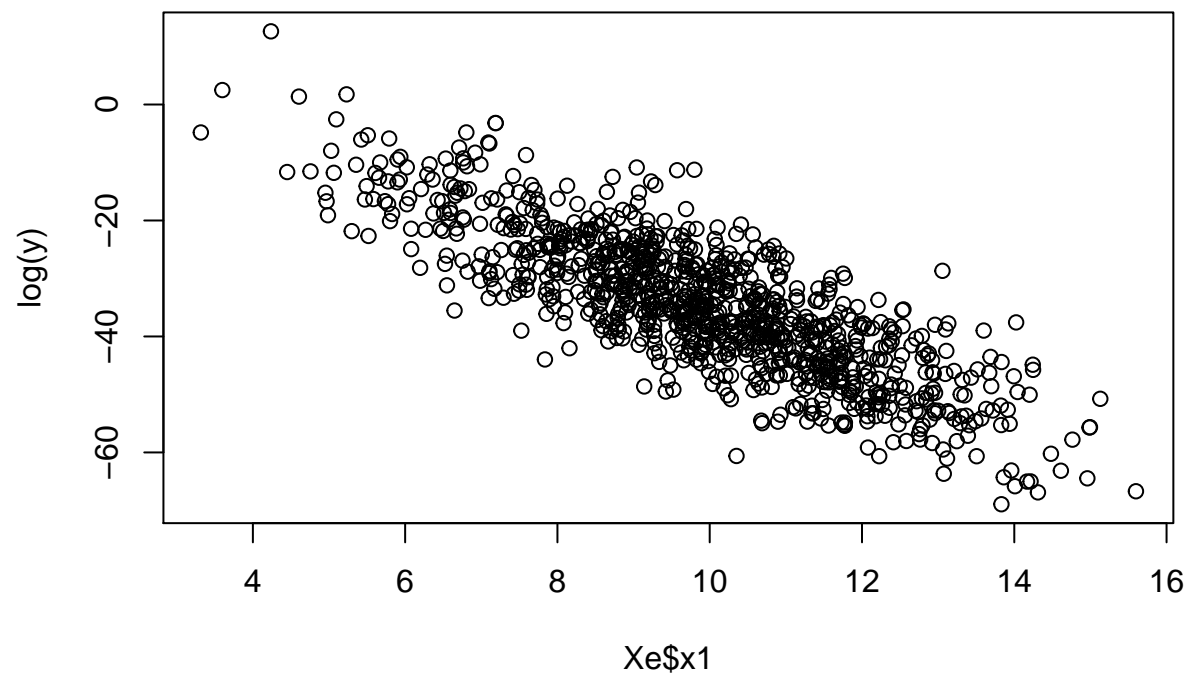


- Generate the true model for outcome y

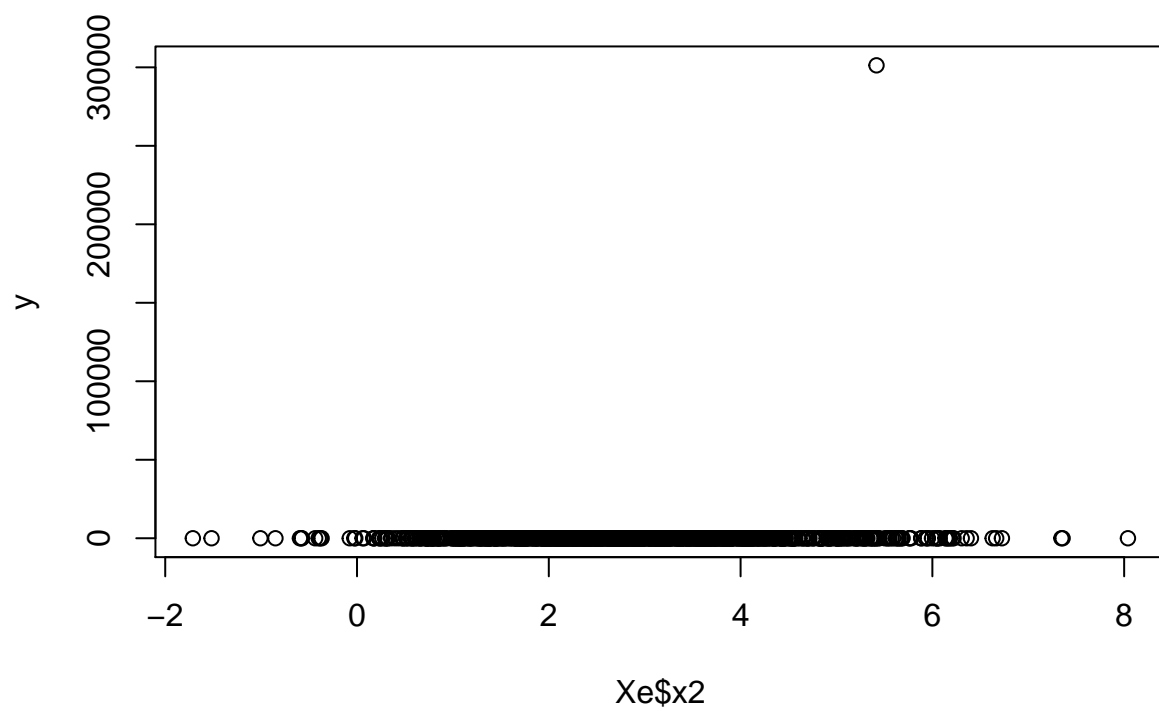
```
y <- exp(10 - 6*Xe$x1 + 5*Xe$x2 + Xe$e) # log(y) is linear in parameters and
                                         # residual, but y is not.
plot(Xe$x1,y)
```



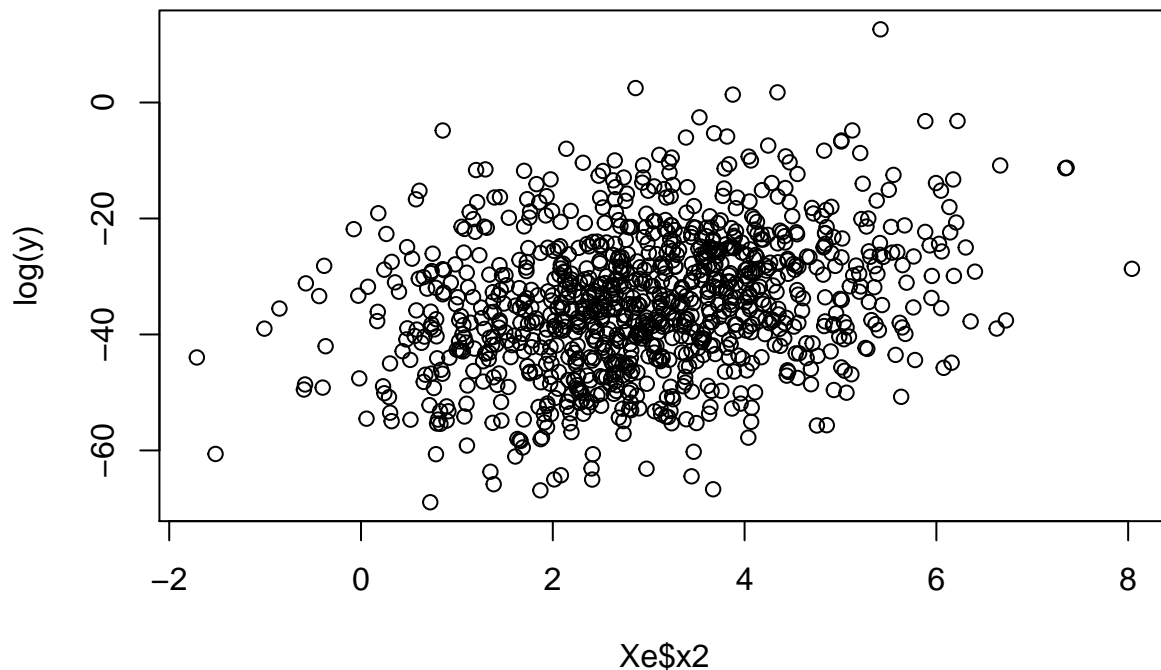
```
plot(Xe$x1,log(y))
```



```
plot(Xe$x2,y)
```

```
plot(Xe$x2,log(y))
```



- Run the linear regression when the true model is linear vs not linear.

```
summary(lm(y~x1+x2,data=Xe))
```

```
##
## Call:
## lm(formula = y ~ x1 + x2, data = Xe)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -3802  -1104   -264    495  296312
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   4009.1     1519.8   2.638 0.008470 **
## x1            -557.4     156.9  -3.553 0.000399 ***
## x2             609.8     222.1   2.746 0.006148 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 9462 on 997 degrees of freedom
## Multiple R-squared:  0.01545,    Adjusted R-squared:  0.01347
## F-statistic: 7.822 on 2 and 997 DF,  p-value: 0.000426
```

```
summary(lm(log(y)~x1+x2,data=Xe))
```

```
##
## Call:
```

```
## lm(formula = log(y) ~ x1 + x2, data = Xe)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -3.3647 -0.6814 -0.0060  0.6925  3.0024
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 10.19399    0.16125   63.22  <2e-16 ***
## x1          -6.02720    0.01665  -362.07  <2e-16 ***
## x2           5.01926    0.02357   212.99  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.004 on 997 degrees of freedom
## Multiple R-squared:  0.9931, Adjusted R-squared:  0.993
## F-statistic: 7.126e+04 on 2 and 997 DF,  p-value: < 2.2e-16
```

Violate Assumption 2: Right hand side variables are constants or are perfectly correlated

- NOTE: The intercept is a constant, but no other variables can be constant
 - otherwise perfectly correlated with the intercept
- Set the seed for replicability.
- Generate the true model

```
set.seed(826)
# Covariance matrix of x1 and e
Sig <- matrix(c(4,0,0,1),2,2)
Sig      # Notice e does not covary with x1 (assumption 5)

##      [,1] [,2]
## [1,]    4    0
## [2,]    0    1

# We will make x2 an exact multiple of x1 (violate assumption 2)
# Mean of x1 and e
moo <- c(10,0)
# generate data
Xe <- mvrnorm(n=1000,mu=moo,Sigma=Sig)
# generate x2 as arbitrary multiple of x1
x2 <- 7*Xe[,1]
Xe <- as.data.frame(cbind(Xe[,1],x2,Xe[,2]))
# give the variables names
colnames(Xe)<-c("x1","x2","e")
head(Xe)

##      x1      x2      e
## 1 10.22234 71.55639 -0.06505297
## 2 11.68949 81.82641 -0.17189937
## 3 10.21010 71.47071  0.36952559
## 4 13.64180 95.49257  0.19523319
## 5 10.76801 75.37608 -1.73650622
## 6 14.34835 100.43846 -0.10192330
```

- Investigate sample correlations and covariances
 - notice difference from “truth”

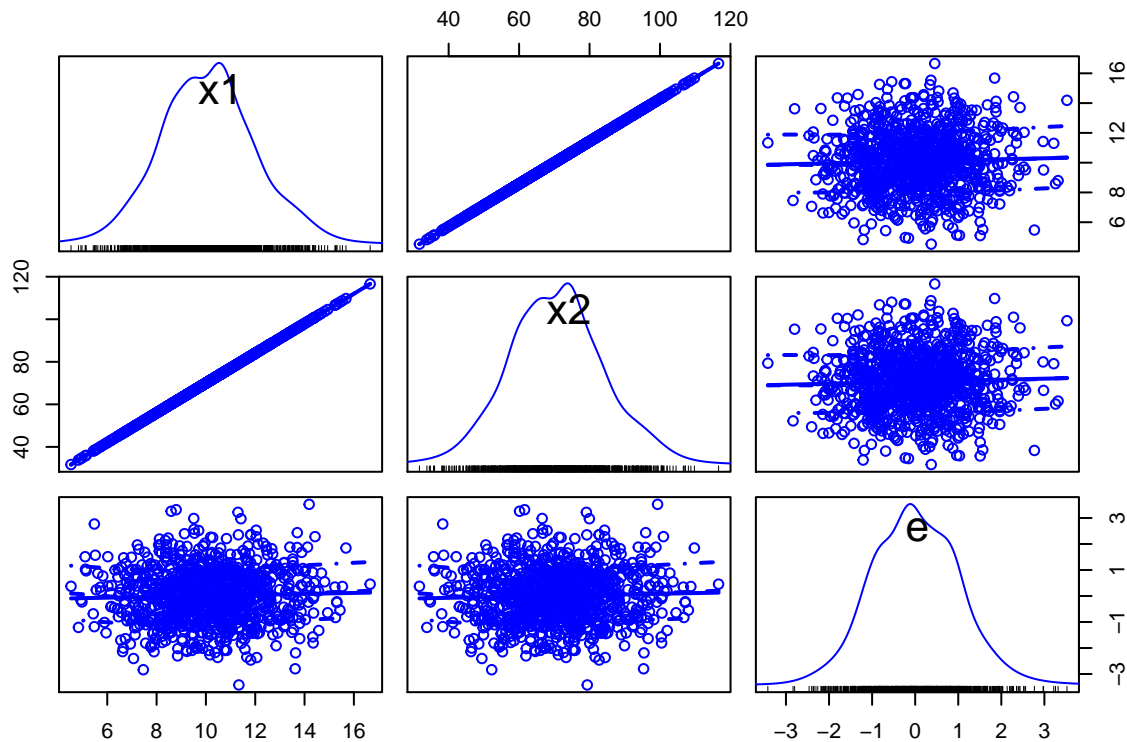
```
cov(Xe)
```

```
##           x1           x2           e
## x1  3.93549450  27.5484615  0.07180935
## x2  27.54846153 192.8392307  0.50266547
## e   0.07180935   0.5026655  1.06201033
```

```
cor(Xe)
```

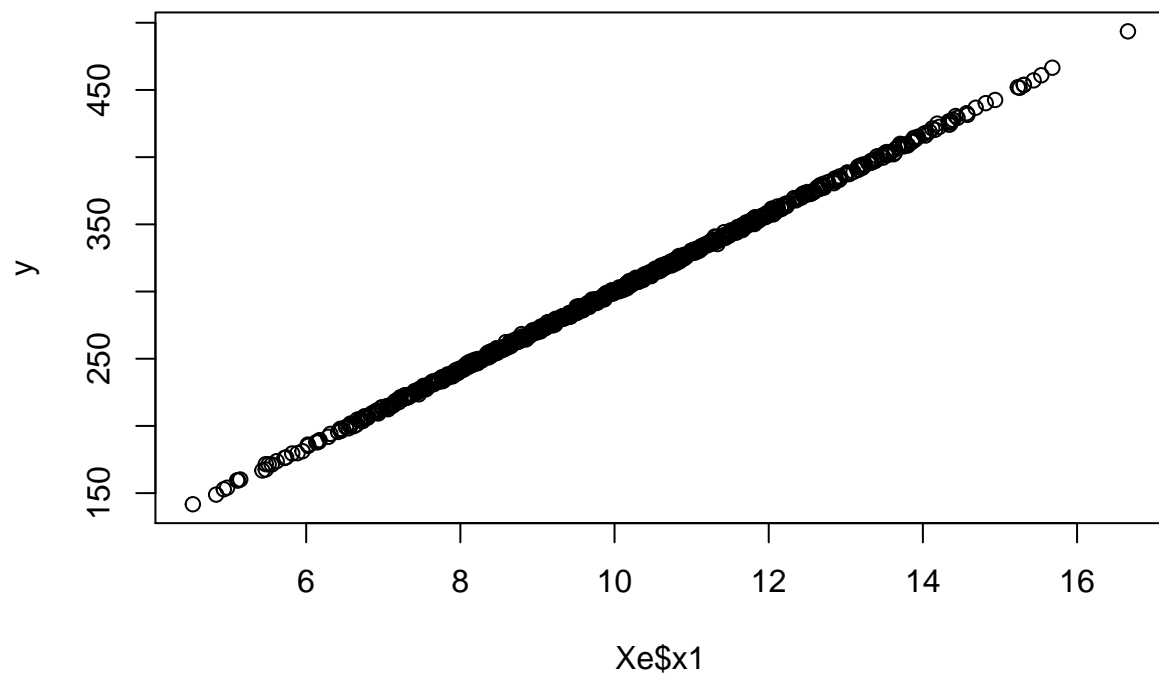
```
##           x1           x2           e
## x1  1.00000000  1.00000000  0.03512505
## x2  1.00000000  1.00000000  0.03512505
## e   0.03512505  0.03512505  1.00000000
```

```
scatterplotMatrix(Xe)
```

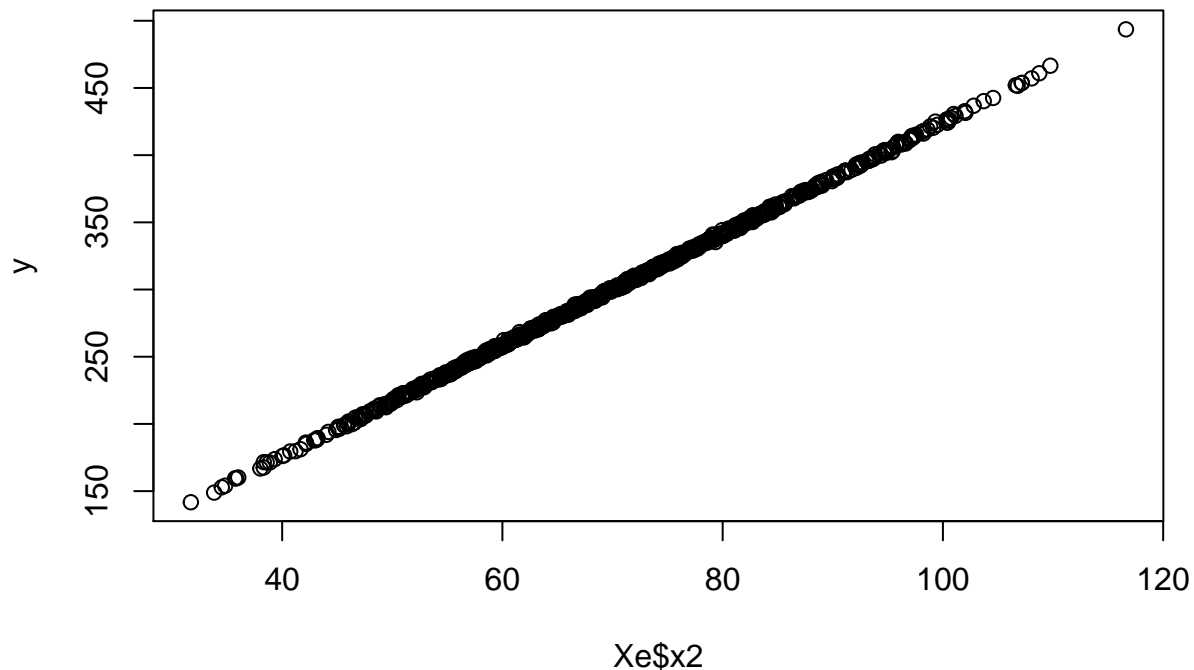


- Generate the true y (outcome)

```
y <- 10 - 6*Xe$x1 + 5*Xe$x2 + Xe$e # y is linear in parameters (assumption 1)
plot(Xe$x1,y)
```



```
plot(Xe$x2,y) # Should be a tight fit!
```



- Run the linear regression

```
summary(lm(y~x1+x2,data=Xe)) # Why an intercept of 10 and coefficient of 29?
```

```
##
## Call:
## lm(formula = y ~ x1 + x2, data = Xe)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -3.4532 -0.7331 -0.0450  0.7074  3.4402
##
## Coefficients: (1 not defined because of singularities)
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  9.82550    0.16904   58.12  <2e-16 ***
## x1          29.01825    0.01643 1765.79  <2e-16 ***
## x2              NA           NA      NA      NA
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.03 on 998 degrees of freedom
## Multiple R-squared:  0.9997, Adjusted R-squared:  0.9997
## F-statistic: 3.118e+06 on 1 and 998 DF, p-value: < 2.2e-16
```

```
# -6 + 5*7 = 29
# some stats packages will not produce output
```

```
# NOTE: regressing on an intercept ONLY estimates the mean of y - intercept can  
# be constant.
```

```
mean(y)
```

```
## [1] 302.725
```

```
summary(lm(y~1))
```

```
##
```

```
## Call:
```

```
## lm(formula = y ~ 1)
```

```
##
```

```
## Residuals:
```

```
##      Min       1Q   Median       3Q      Max  
## -160.99  -39.38   -0.12   35.42  190.88
```

```
##
```

```
## Coefficients:
```

```
##              Estimate Std. Error t value Pr(>|t|)  
## (Intercept)  302.725      1.821   166.3  <2e-16 ***
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##
```

```
## Residual standard error: 57.58 on 999 degrees of freedom
```

- But what if another regressor is also constant?

```
set.seed(826)
```

```
# Covariance matrix of x1, x2, and e
```

```
Sig <- matrix(c(4,0,0,0,0,0,1),3,3)
```

```
Sig      # Notice e does not covary with x1 or x2 (assumption 5)
```

```
##      [,1] [,2] [,3]
```

```
## [1,]    4    0    0
```

```
## [2,]    0    0    0
```

```
## [3,]    0    0    1
```

```
# But x2 also has no variance (violate assumption 2)
```

```
# Mean of x1, x2, and e
```

```
moo <- c(10,3,0)
```

```
# generate data
```

```
Xe <- mvrnorm(n=1000,mu=moo,Sigma=Sig)
```

```
# give the variables names
```

```
colnames(Xe)<-c("x1","x2","e")
```

```
Xe <- as.data.frame(Xe)
```

```
head(Xe)
```

```
##      x1 x2      e
```

```
## 1 9.777659 3 0.06505297
```

```
## 2 8.310513 3 0.17189937
```

```
## 3 9.789899 3 -0.36952559
```

```
## 4 6.358205 3 -0.19523319
```

```
## 5 9.231989 3 1.73650622
```

```
## 6 5.651648 3 0.10192330
```

```
# Sample correlations and covariances (notice difference from "truth")
```

```
cov(Xe)
```

```
##           x1 x2          e
## x1 3.93549450 0 0.07180935
## x2 0.00000000 0 0.00000000
## e  0.07180935 0 1.06201033
```

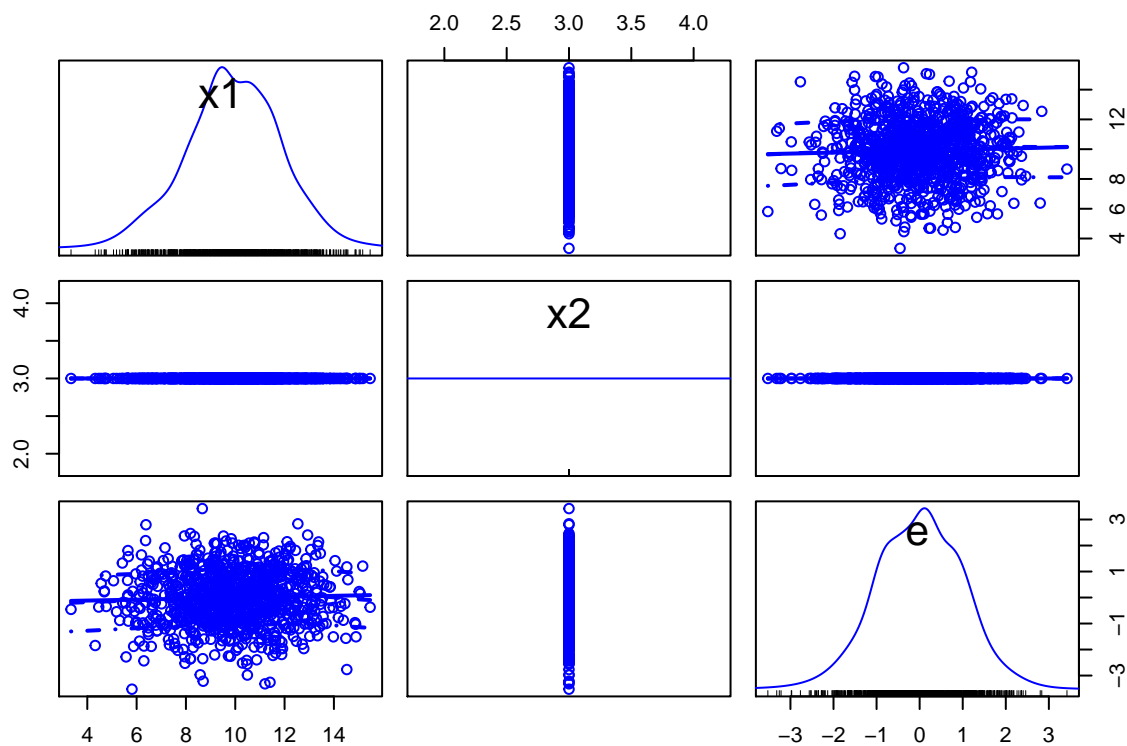
```
cor(Xe)
```

```
## Warning in cor(Xe): the standard deviation is zero
```

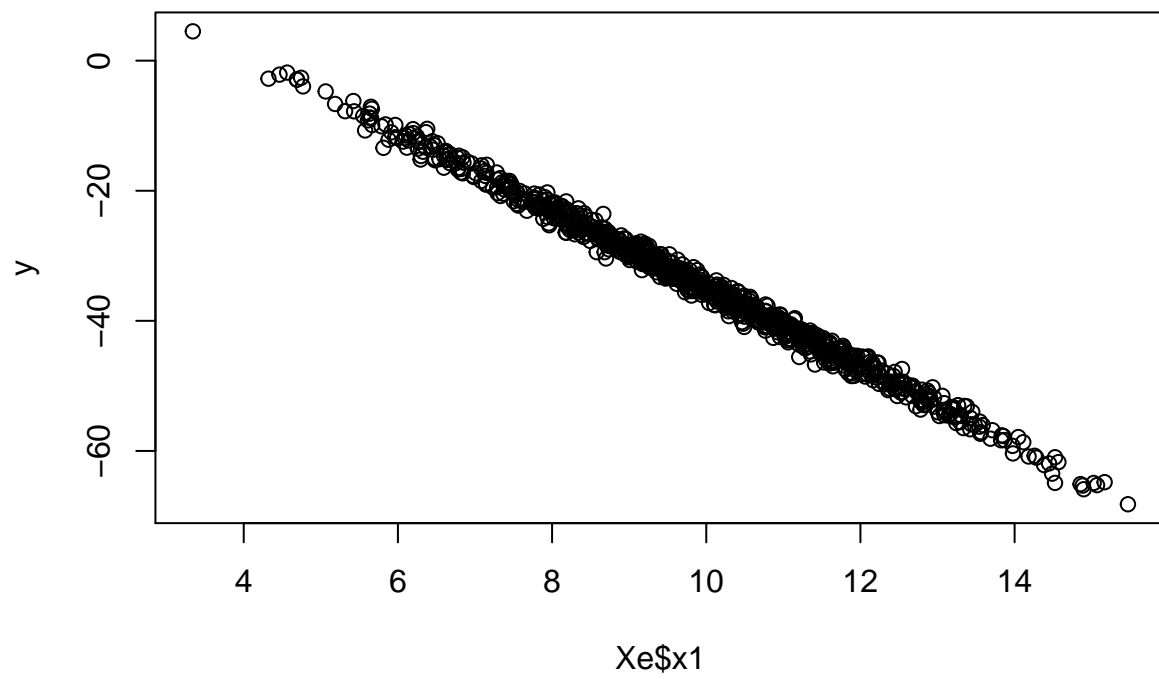
```
##           x1 x2          e
## x1 1.00000000 NA 0.03512505
## x2          NA  1          NA
## e  0.03512505 NA 1.00000000
```

```
scatterplotMatrix(Xe)
```

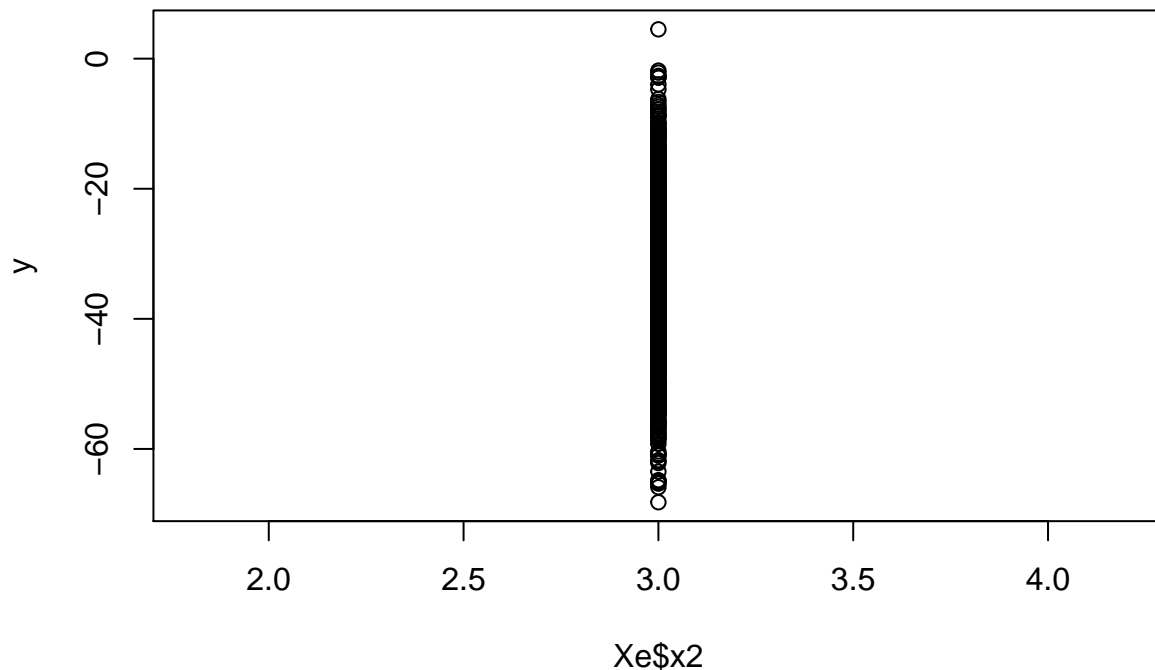
```
## Warning in smoother(x[subs], y[subs], col = smoother.args$col[i], log.x =
## FALSE, : could not fit negative part of the spread
```



```
# generate the true y (outcome)
y <- 10 - 6*Xe$x1 + 5*Xe$x2 + Xe$e # y is linear in parameters (assumption 1)
plot(Xe$x1,y)
```

```
plot(Xe$x2,y) # Should be a tight fit!
```



```
# run the linear regression
summary(lm(y~x1+x2,data=Xe)) # Why an intercept of 25 and coefficient of -6?
```

```
##
## Call:
## lm(formula = y ~ x1 + x2, data = Xe)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -3.4402 -0.7074  0.0450  0.7331  3.4532
##
## Coefficients: (1 not defined because of singularities)
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  24.80957    0.16603   149.4   <2e-16 ***
## x1          -5.98175    0.01643  -364.0   <2e-16 ***
## x2                   NA           NA      NA      NA
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.03 on 998 degrees of freedom
## Multiple R-squared:  0.9925, Adjusted R-squared:  0.9925
## F-statistic: 1.325e+05 on 1 and 998 DF,  p-value: < 2.2e-16
```

```
# 10 + 5*3
```

```
# some stats packages will not even produce output
```

Violate Assumption 3: Residuals “e” do not have constant variance

- Residuals are “heteroskedastic” (different variance)
 - Variance changes at different places in the population
 - For different values of X or different time periods t
- Set the seed and generate the true model

```
set.seed(826)
# Covariance matrix of x1, x2
Sig <- matrix(c(4,1,1,2),2,2)
Sig

##      [,1] [,2]
## [1,]    4    1
## [2,]    1    2
# Mean of x1, x2
moo <- c(10,3)
# generate data
Xe <- mvrnorm(n=1000,mu=moo,Sigma=Sig)

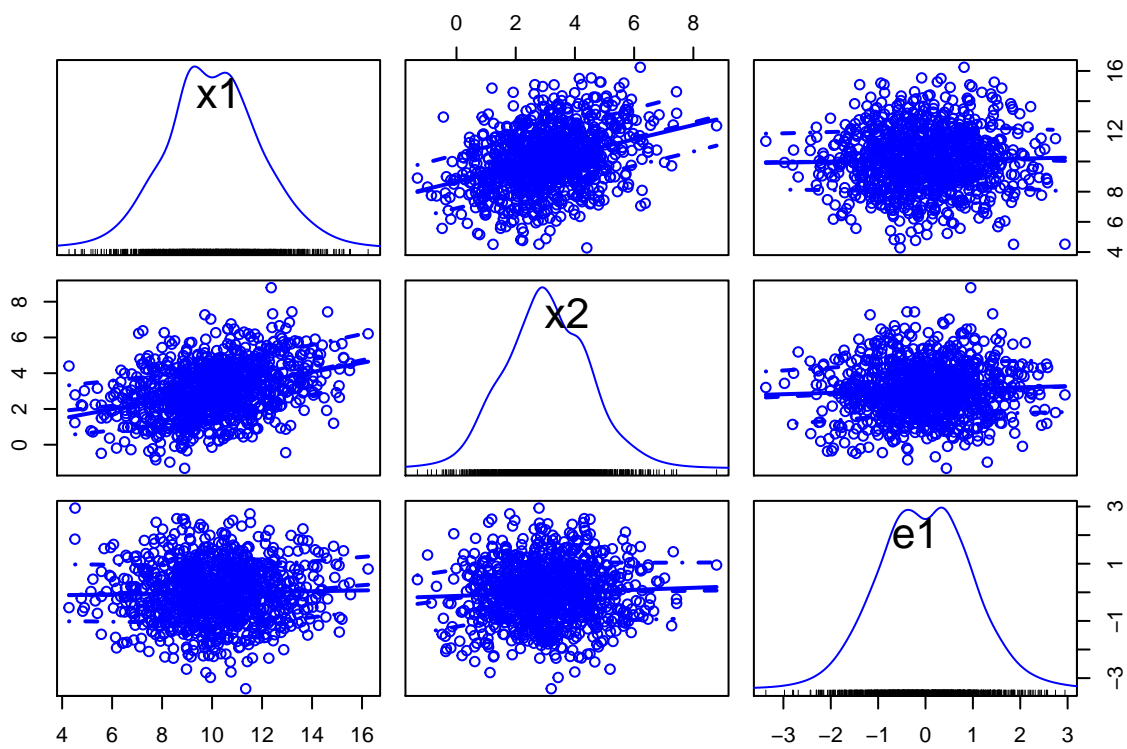
# generate homoskedastic residuals
eps = rnorm(n=1000,mean=0,sd=sqrt(1))

# generate heteroskedastic residuals
sigma2 = (eps^2)*(Xe[,1]^2+Xe[,2]^2)
eps2 = rnorm(n=1000,mean=0,sd=sqrt(sigma2))

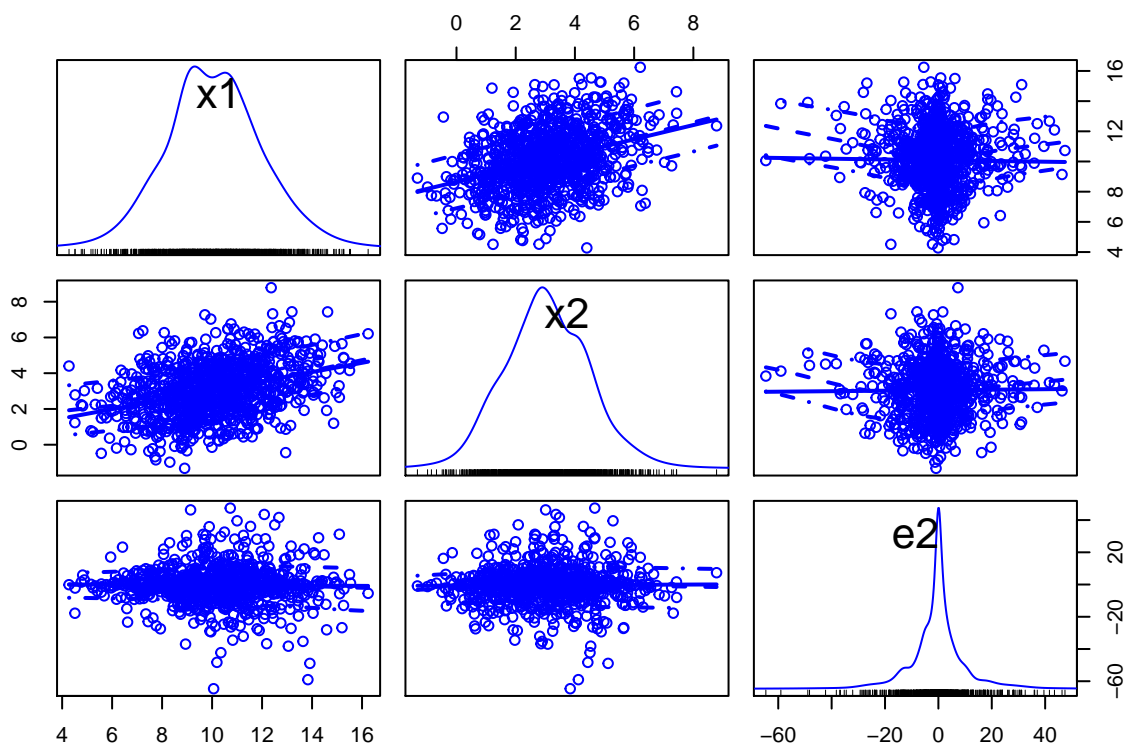
Xe1 = as.data.frame(cbind(Xe,eps))
colnames(Xe1)<-c("x1", "x2", "e1")
Xe2 = as.data.frame(cbind(Xe,eps2))
colnames(Xe2)<-c("x1", "x2", "e2")
```

- Investigate sample correlations and covariances
 - notice difference from “truth”

```
scatterplotMatrix(Xe1)
```



```
scatterplotMatrix(Xe2)
```

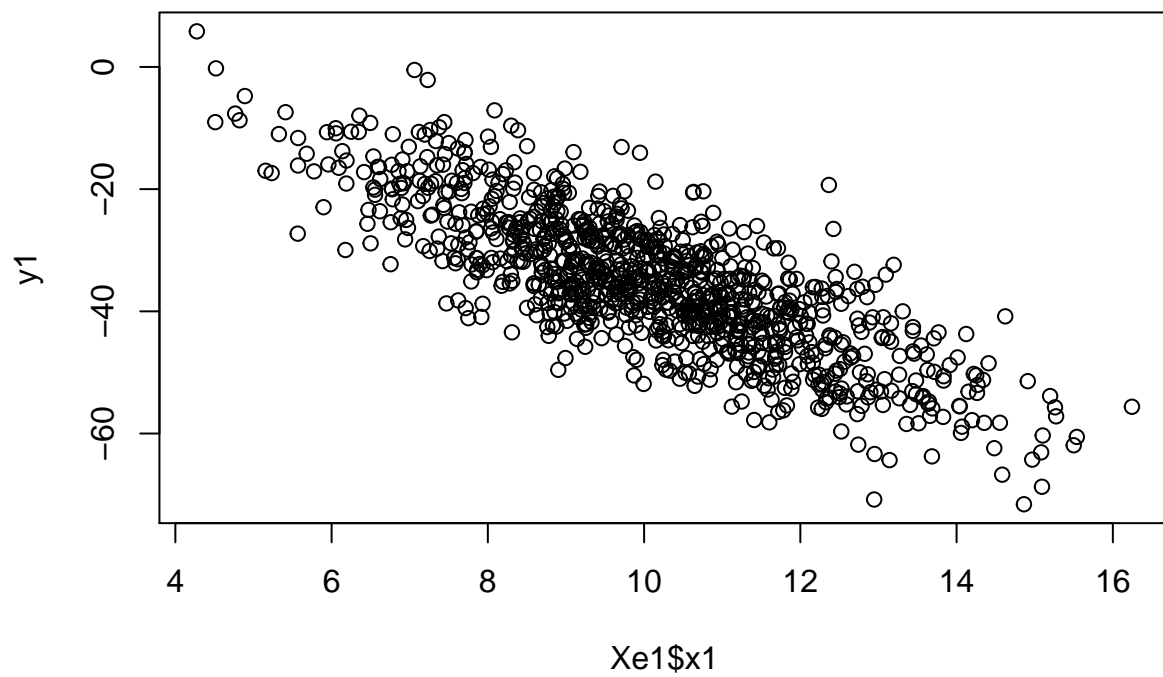


- Generate the true y (outcome)

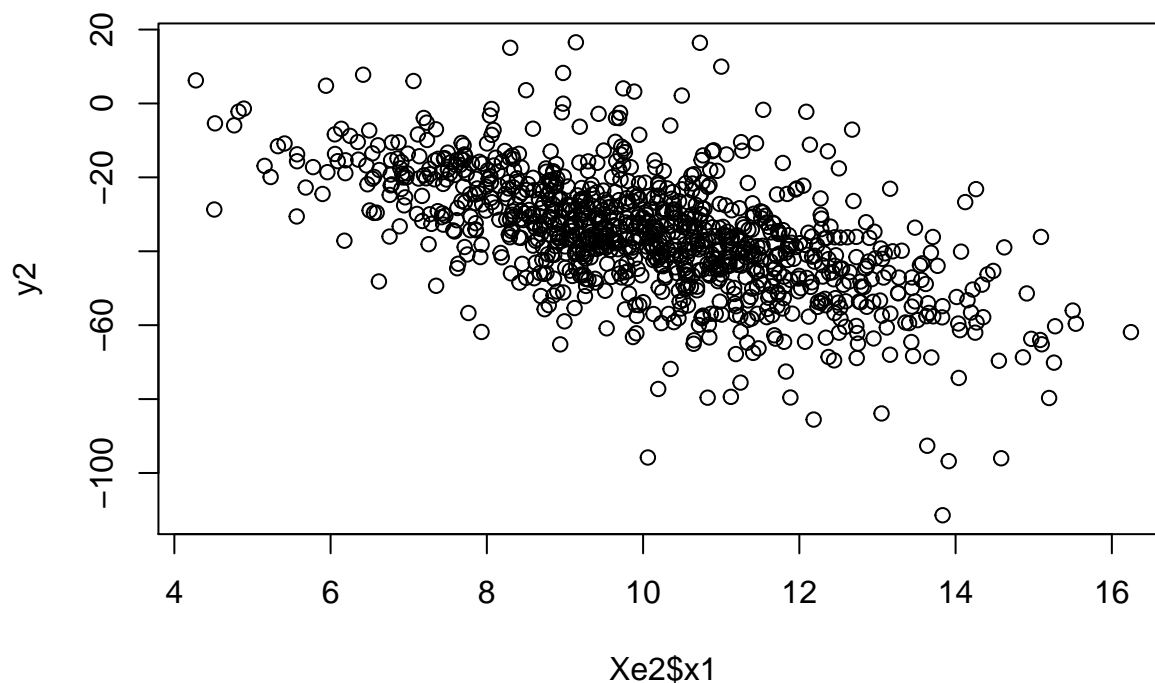
```
y1 <- 10 - 6*Xe1$x1 + 5*Xe1$x2 + Xe1$e1
y2 <- 10 - 6*Xe2$x1 + 5*Xe2$x2 + Xe2$e2
```

- Notice there are not necessarily obvious differences in the plot

```
plot(Xe1$x1,y1)
```



```
plot(Xe2$x1,y2)
```



- Run the linear regression
 - notice the difference in residual standard error, coefficient std. error, Rsquared.

```
summary(lm(y1~x1+x2,data=Xe1))
```

```
##
## Call:
## lm(formula = y1 ~ x1 + x2, data = Xe1)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -3.3647 -0.6814 -0.0060  0.6925  3.0024
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   9.82512    0.16589   59.23  <2e-16 ***
## x1            -5.99439    0.01719  -348.62  <2e-16 ***
## x2             5.03285    0.02317   217.23  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.004 on 997 degrees of freedom
## Multiple R-squared:  0.9925, Adjusted R-squared:  0.9925
## F-statistic: 6.596e+04 on 2 and 997 DF,  p-value: < 2.2e-16
summary(lm(y2~x1+x2,data=Xe2))
```

```
##
## Call:
## lm(formula = y2 ~ x1 + x2, data = Xe2)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -64.361  -4.232   0.309   3.882  47.679
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  10.3175     1.8010   5.729 1.34e-08 ***
## x1           -6.1093     0.1867 -32.726 < 2e-16 ***
## x2            5.1291     0.2515  20.391 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 10.9 on 997 degrees of freedom
## Multiple R-squared:  0.5383, Adjusted R-squared:  0.5374
## F-statistic: 581.2 on 2 and 997 DF, p-value: < 2.2e-16
```

Violate Assumption 4: Residuals “e” are serially correlated (autocorrelated)

- Set the seed and generate the true model

```
set.seed(826)
# Covariance matrix of x1, x2
Sig <- matrix(c(4,1,1,2),2,2)
Sig

##      [,1] [,2]
## [1,]    4    1
## [2,]    1    2

# Mean of x1, x2
moo <- c(10,3)
# generate data
Xe <- mvrnorm(n=1000,mu=moo,Sigma=Sig)

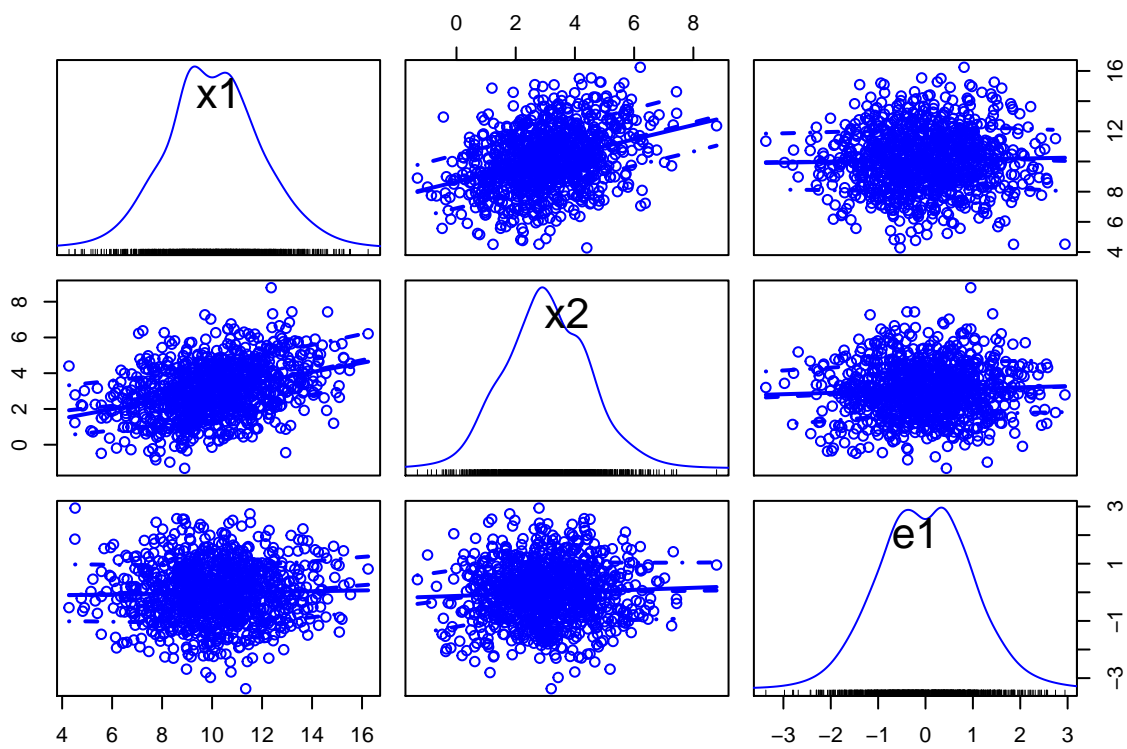
# generate independent residuals
eps = rnorm(n=1000,mean=0,sd=sqrt(1))

# generate serially correlated residuals
eps2 <- arima.sim(model=list(ar=c(0.8)),n=1000,sd=1)

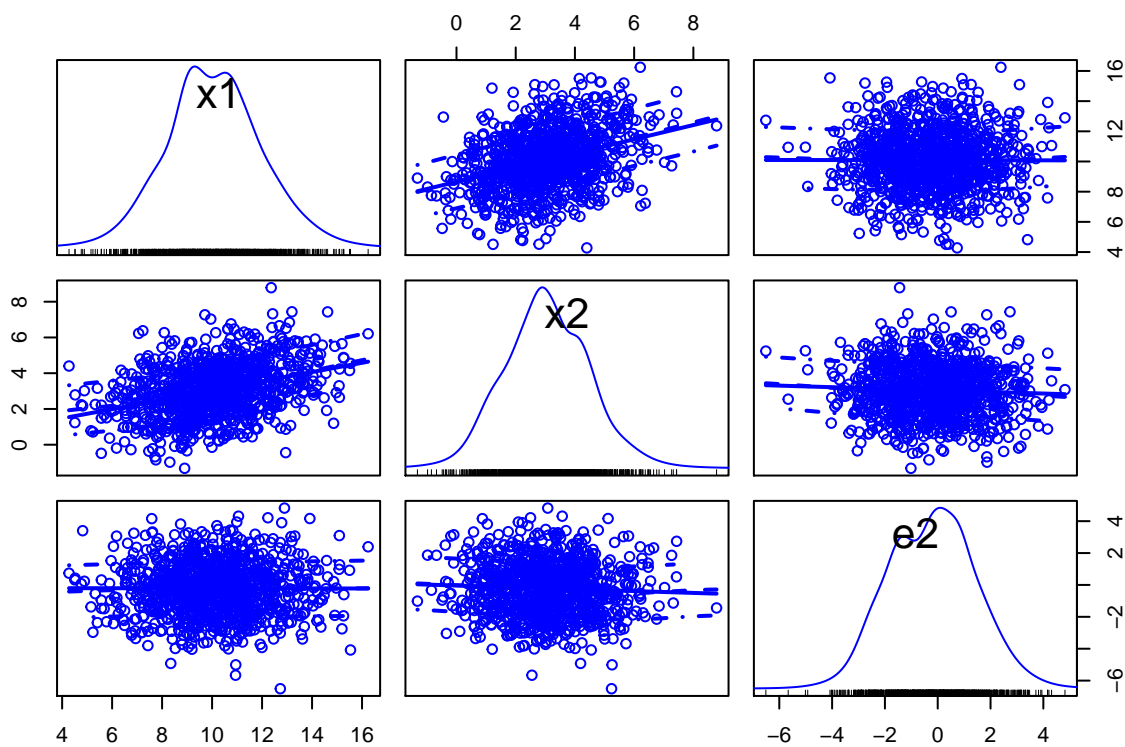
Xe1 = as.data.frame(cbind(Xe,eps))
colnames(Xe1)<-c("x1","x2","e1")
Xe2 = as.data.frame(cbind(Xe,eps2))
colnames(Xe2)<-c("x1","x2","e2")
```

- Investigate sample correlations and covariances
 - notice difference from “truth”

```
scatterplotMatrix(Xe1)
```

```
scatterplotMatrix(Xe2)
```

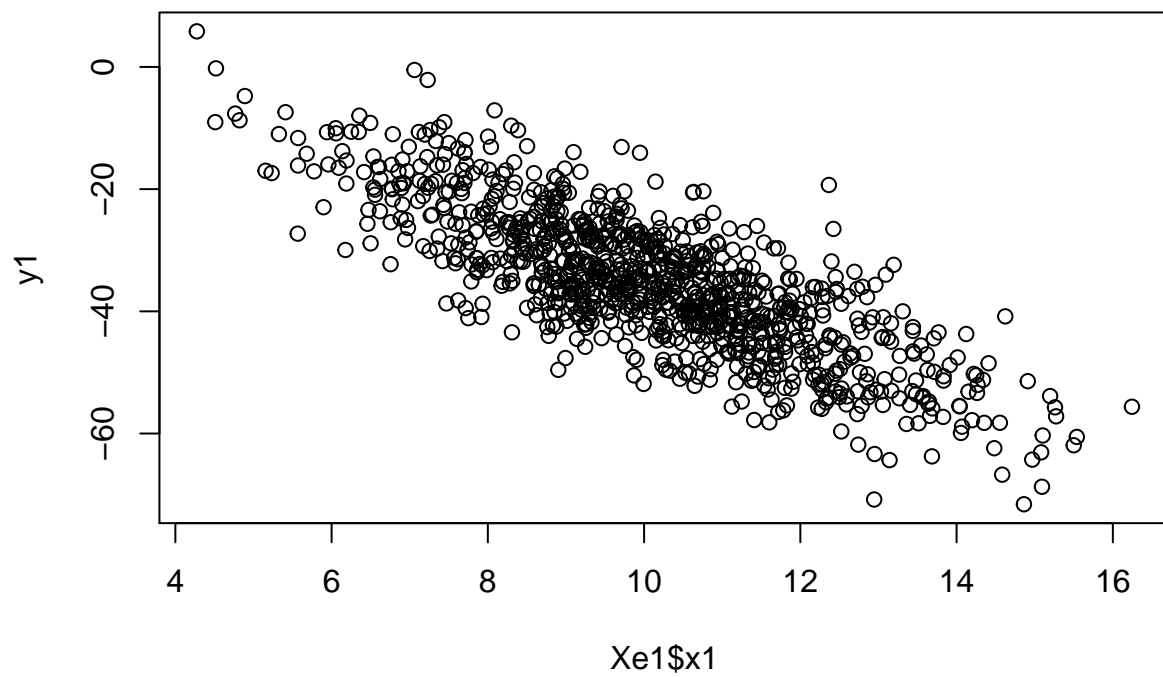


- Generate the true y (outcome)

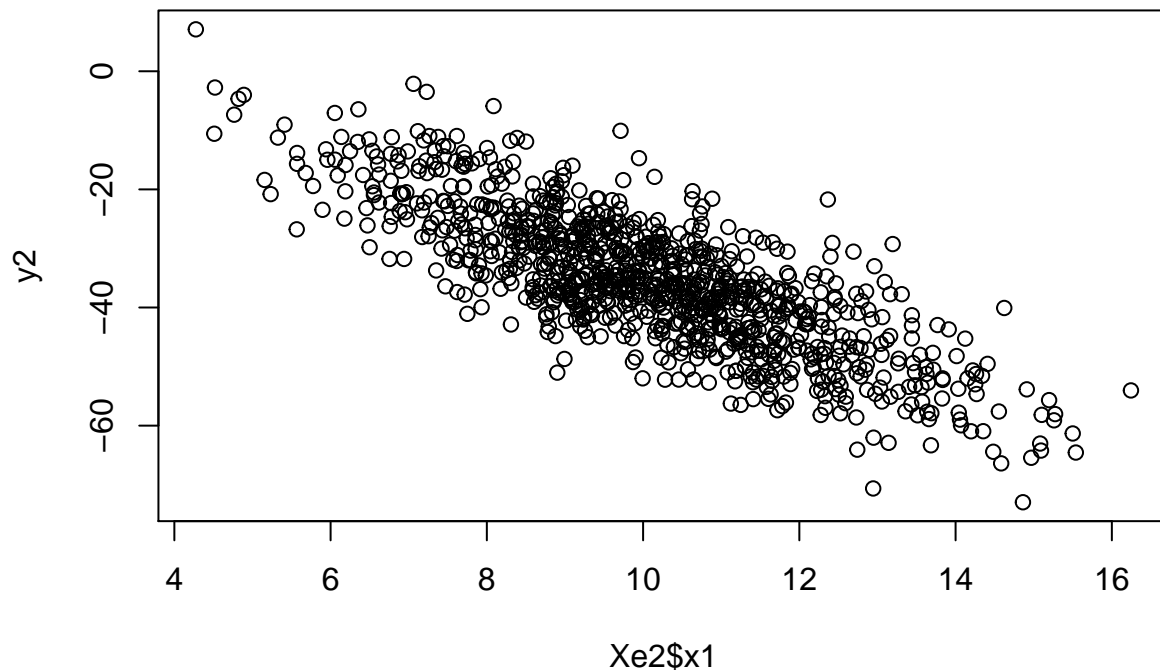
```
y1 <- 10 - 6*Xe1$x1 + 5*Xe1$x2 + Xe1$e1
y2 <- 10 - 6*Xe2$x1 + 5*Xe2$x2 + Xe2$e2
```

- Note that there are not necessarily obvious differences in the plot

```
plot(Xe1$x1,y1)
```



```
plot(Xe2$x1,y2)
```



- Run the linear regression
 - notice the difference in residual standard error, coefficient std. error, Rsquared.

```
summary(lm(y1~x1+x2,data=Xe1))
```

```
##
## Call:
## lm(formula = y1 ~ x1 + x2, data = Xe1)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -3.3647 -0.6814 -0.0060  0.6925  3.0024
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   9.82512    0.16589   59.23  <2e-16 ***
## x1            -5.99439    0.01719  -348.62  <2e-16 ***
## x2             5.03285    0.02317   217.23  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.004 on 997 degrees of freedom
## Multiple R-squared:  0.9925, Adjusted R-squared:  0.9925
## F-statistic: 6.596e+04 on 2 and 997 DF,  p-value: < 2.2e-16
```

```
summary(lm(y2~x1+x2,data=Xe2))
```

```
##
## Call:
## lm(formula = y2 ~ x1 + x2, data = Xe2)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -6.1938 -1.2369  0.0212  1.1345  4.9921
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   9.83236    0.27727   35.46  <2e-16 ***
## x1           -5.98476    0.02874 -208.24  <2e-16 ***
## x2            4.93465    0.03872  127.43  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.678 on 997 degrees of freedom
## Multiple R-squared:  0.9791, Adjusted R-squared:  0.9791
## F-statistic: 2.339e+04 on 2 and 997 DF,  p-value: < 2.2e-16
```

Violate Assumption 5: Residuals “e” are correlated with X variables

- Some X variables are **endogenous**
- Many flavors of this

```
set.seed(826)
# Covariance matrix of x1, x2, and e
n <- 3
A <- matrix(runif(n^2)*2-1, ncol=n)
Sig <- t(A) %*% A
Sig      # Notice e covaries with x1 and x2
```

5a. Residuals (unobservable) covary with X's (observable)

```
##           [,1]      [,2]      [,3]
## [1,]  1.1012487 -0.4083661  0.5785058
## [2,] -0.4083661  1.4481911  0.6653081
## [3,]  0.5785058  0.6653081  1.0935305

# also x1 and x2 can covary, but not perfectly (assumption 2)
# Mean of x1, x2, and e
moo <- c(10,3,0)
# generate data
Xe1 <- mvrnorm(n=1000,mu=moo,Sigma=Sig)

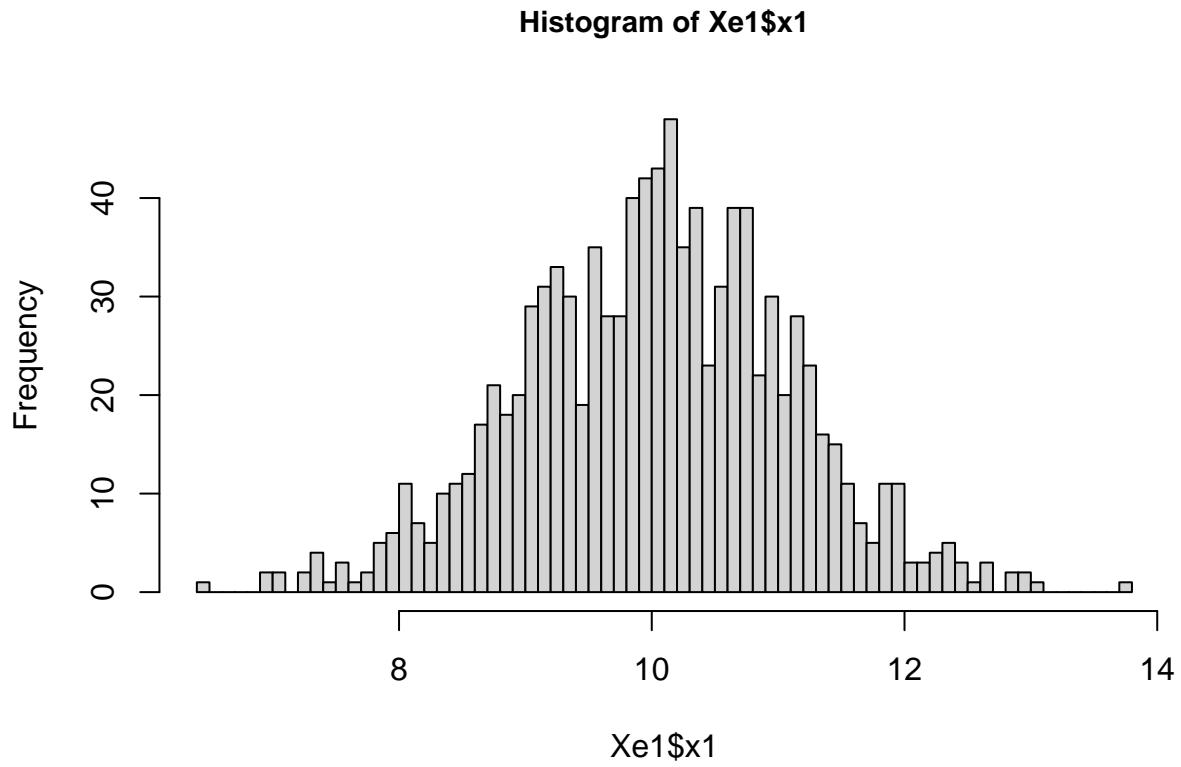
# give the variables names
colnames(Xe1)<-c("x1","x2","e1")
# store as a data frame
Xe1 <- as.data.frame(Xe1)
head(Xe1)

##           x1           x2           e1
## 1  7.733325  3.255037 -1.6600797
```

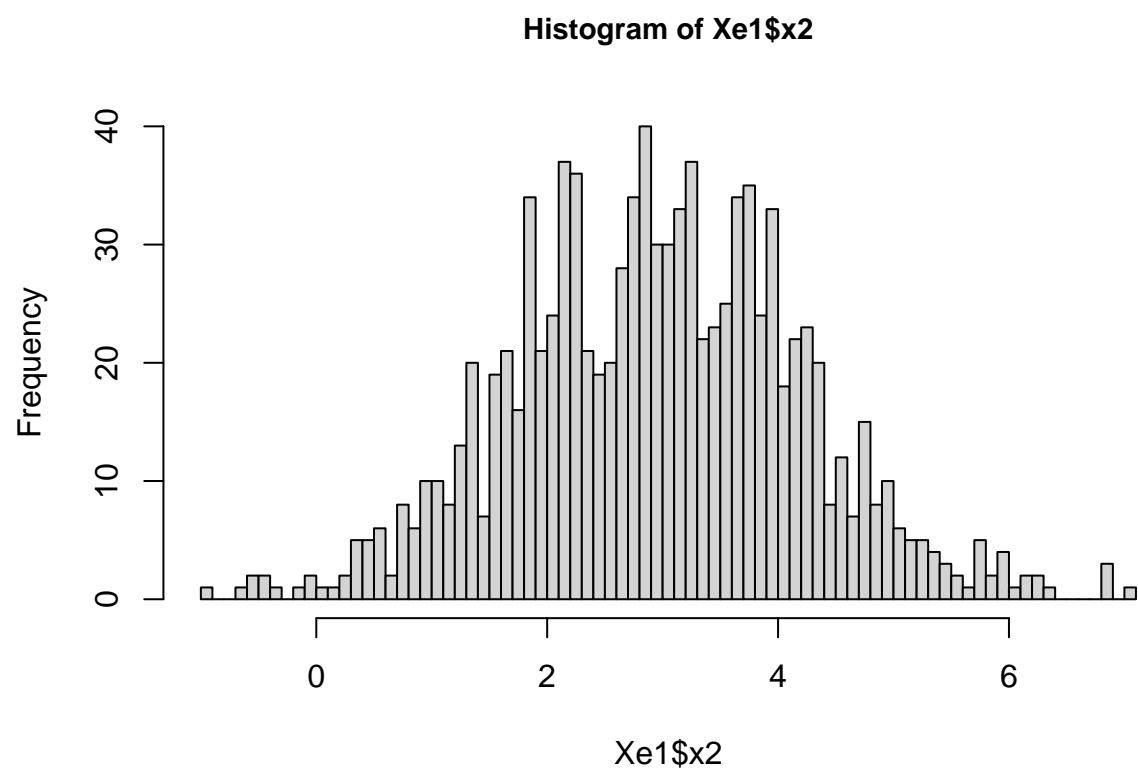
```
## 2 10.031196 4.345290 0.6720729
## 3 10.853379 3.439955 0.6454713
## 4 10.285622 3.138029 0.7302525
## 5 11.973566 3.084988 1.9555648
## 6 9.263035 5.055792 0.2941055
```

- Investigate patterns in the data

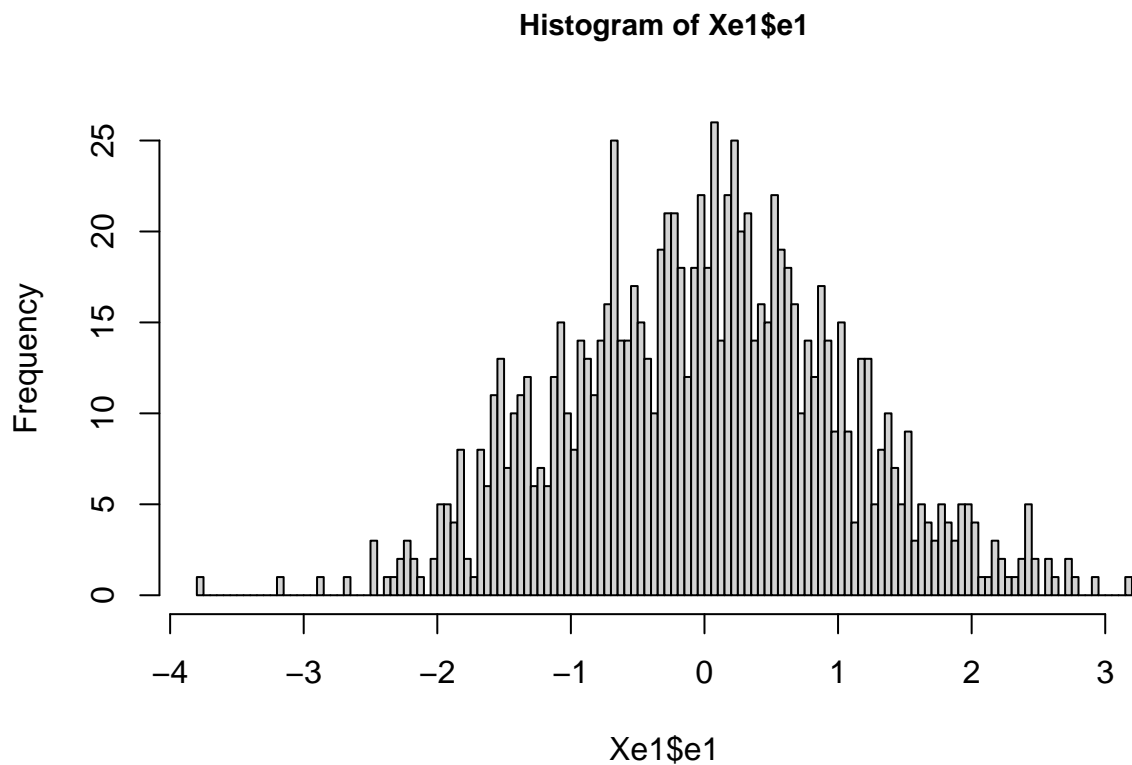
```
# plot empirical distribution of each:
hist(Xe1$x1, breaks = 100, cex.main = 0.9)
```



```
hist(Xe1$x2, breaks = 100, cex.main = 0.9)
```



```
hist(Xe1$x2, breaks = 100, cex.main = 0.9)
```



Sample correlations and covariances (notice difference from "truth")

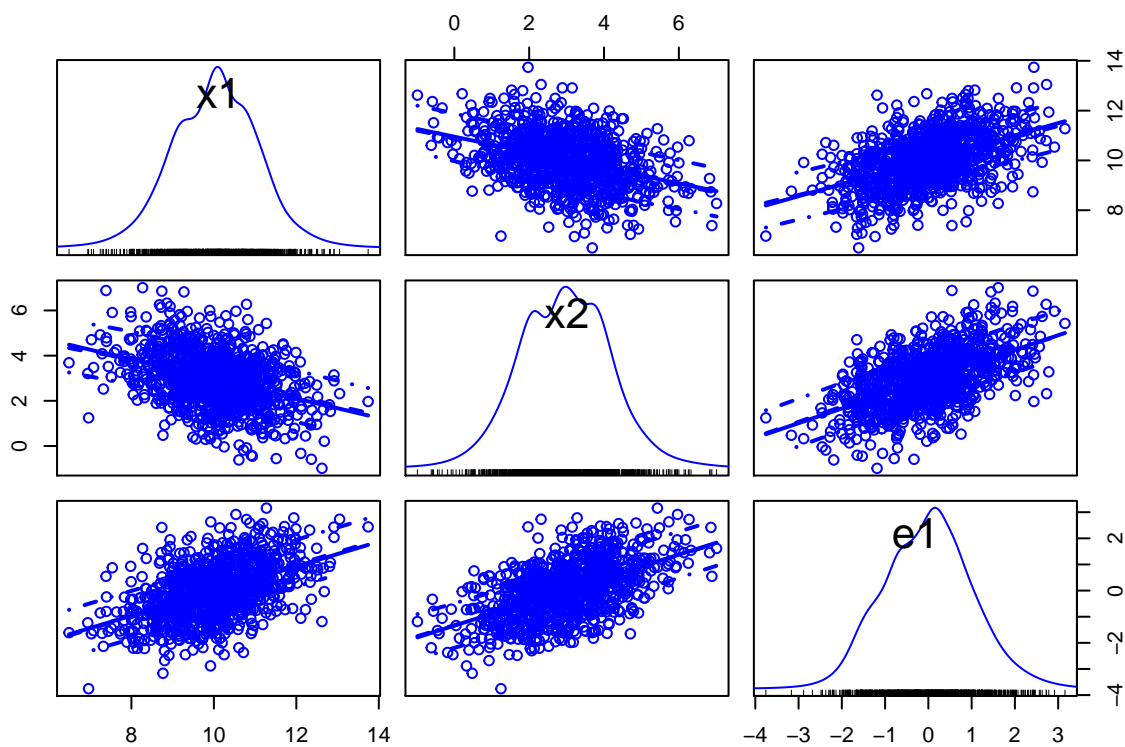
```
cov(Xe1)
```

```
##           x1           x2           e1
## x1  1.1261792 -0.4863156  0.5318302
## x2 -0.4863156  1.5560936  0.7052121
## e1  0.5318302  0.7052121  1.0883354
```

```
cor(Xe1)
```

```
##           x1           x2           e1
## x1  1.0000000 -0.3673640  0.4803833
## x2 -0.3673640  1.0000000  0.5419017
## e1  0.4803833  0.5419017  1.0000000
```

```
scatterplotMatrix(Xe1)
```

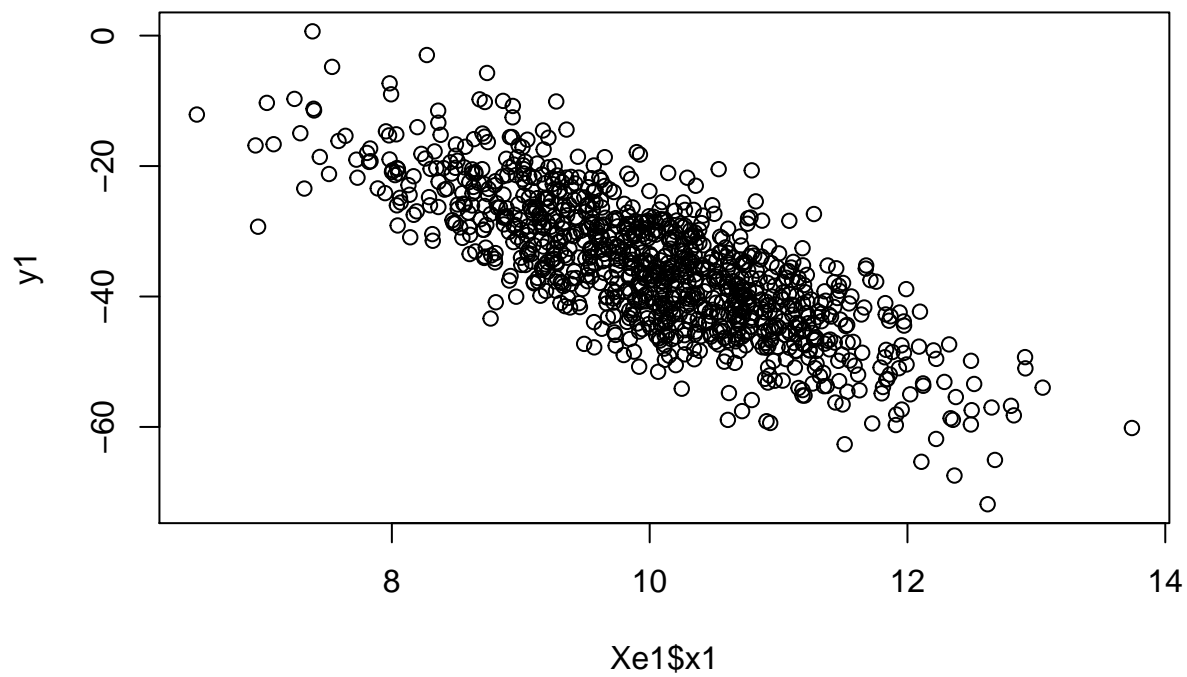



- Generate the true y (outcome)

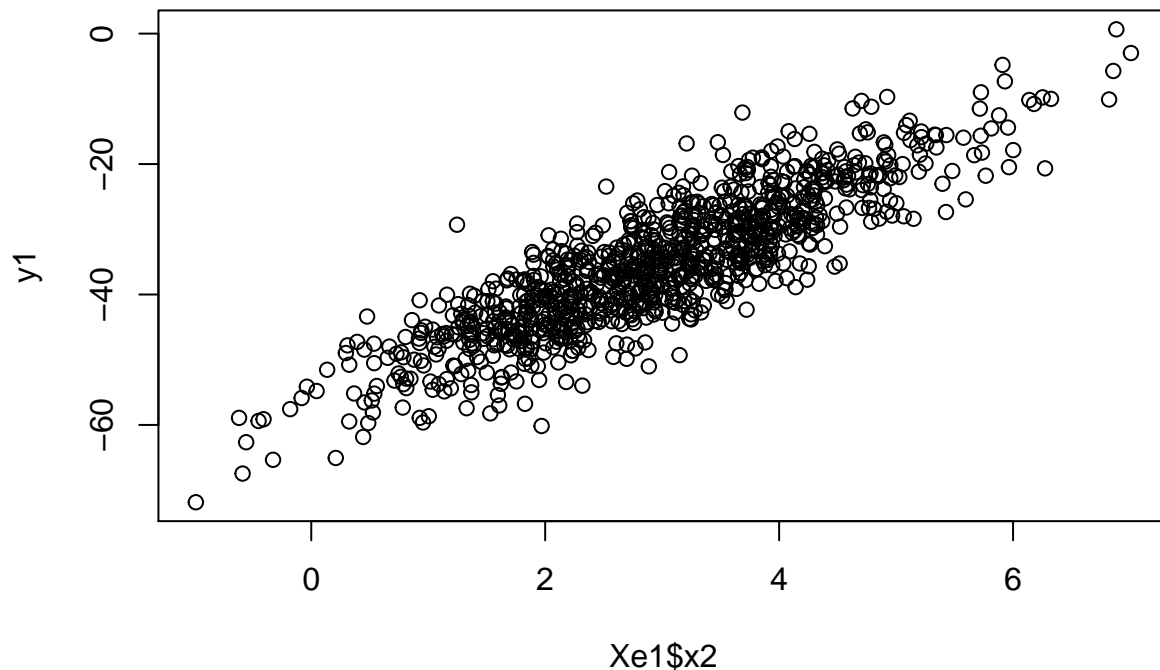
```
y1 <- 10 - 6*Xe1$x1 + 5*Xe1$x2 + Xe1$e1
```

- No obvious problems in the plot of X 's against y

```
plot(Xe1$x1,y1)
```



```
plot(Xe1$x2,y1)
```



- Run the linear regression
 - notice the difference in sample coefficients from true values.

```
summary(lm(y1~x1+x2,data=Xe1))
```

```
##
## Call:
## lm(formula = y1 ~ x1 + x2, data = Xe1)
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-1.54018	-0.29630	-0.01474	0.28478	1.25077

```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.20152	0.15636	1.289	0.198
x1	-5.22785	0.01391	-375.858	<2e-16 ***
x2	5.69451	0.01183	481.250	<2e-16 ***

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.4339 on 997 degrees of freedom
## Multiple R-squared:  0.9983, Adjusted R-squared:  0.9983
## F-statistic: 2.923e+05 on 2 and 997 DF, p-value: < 2.2e-16
```

```
set.seed(826)
# Covariance matrix of x1, x2, and e
Sig <- matrix(c(4,1,0,1,2,0,0,0,1),3,3)
Sig      # Notice e does not covary with x1 or x2 (assumption 5)
```

5b. An observable variable (that is correlated with included variables) was omitted

```
##      [,1] [,2] [,3]
## [1,]    4    1    0
## [2,]    1    2    0
## [3,]    0    0    1

# also x1 and x2 can covary, but not perfectly (assumption 2)
# Mean of x1, x2, and e
moo <- c(10,3,0)

# generate data
Xe <- mvrnorm(n=1000,mu=moo,Sigma=Sig)
# give the variables names
colnames(Xe)<-c("x1","x2","e")
# store as a data frame
Xe <- as.data.frame(Xe)
```

- Generate the true y (outcome)

```
y2 <- 10 - 6*Xe$x1 + 5*Xe$x2 + Xe$e
```

- Run the linear regression with an omitted variable

```
summary(lm(y2~x1+x2,data=Xe))

##
## Call:
## lm(formula = y2 ~ x1 + x2, data = Xe)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -3.3647 -0.6814 -0.0060  0.6925  3.0024
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  10.19399    0.16125   63.22  <2e-16 ***
## x1          -6.02720    0.01665  -362.07  <2e-16 ***
## x2           5.01926    0.02357   212.99  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.004 on 997 degrees of freedom
## Multiple R-squared:  0.9931, Adjusted R-squared:  0.993
## F-statistic: 7.126e+04 on 2 and 997 DF, p-value: < 2.2e-16

summary(lm(y2~x1,data=Xe))
```

```
##
## Call:
## lm(formula = y2 ~ x1, data = Xe)
```

```
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -23.8197  -4.4152  -0.2228   4.5777  22.8015
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   14.327      1.091   13.13  <2e-16 ***
## x1            -4.938      0.108  -45.73  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 6.843 on 998 degrees of freedom
## Multiple R-squared:  0.677, Adjusted R-squared:  0.6766
## F-statistic: 2091 on 1 and 998 DF, p-value: < 2.2e-16
summary(lm(y2~x2,data=Xe))
```

```
##
## Call:
## lm(formula = y2 ~ x2, data = Xe)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -33.793  -7.934  -0.057   7.530  41.335
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -41.7047      0.8497 -49.083  <2e-16 ***
## x2           2.3970      0.2580   9.291  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 11.55 on 998 degrees of freedom
## Multiple R-squared:  0.07961, Adjusted R-squared:  0.07869
## F-statistic: 86.32 on 1 and 998 DF, p-value: < 2.2e-16
```

```
set.seed(826)
# generate an autocorrelated residual
eps <- arima.sim(model=list(ar=c(0.8)),n=999,sd=1)

# generate an autocorrelated outcome that has "eps" as its residual
y1 <- list()
y10 <- rnorm(n=1,mean=10,sd=1)
y1[[1]] <- y10
for(i in 2:1000) {
  y1[[i]] <- 10 + 0.4*y1[[i-1]] + eps[i]
}
y1 <- unlist(y1)

arima(y1,order=c(1,0,0))
```

5c. y is autocorrelated, may or may not have serial correlation/autocorrelation in the residual.

```
##
## Call:
## arima(x = y1, order = c(1, 0, 0))
##
## Coefficients:
##          ar1  intercept
##          0.9146   16.3248
## s.e.  0.0132    0.3819
##
## sigma^2 estimated as 1.086:  log likelihood = -1459.74,  aic = 2925.48
```