

Multivariate Time Series Analysis

Econometrics

Notes draw heavily from the following textbooks: Time Series Analysis (Hamilton, 1994), Econometric Analysis (Greene, 7th edition), and Analysis of Financial Time Series (Tsay, 3rd edition)

F-test review

- Suppose we want to run the regression

$$y_t = \beta_0 + \beta_1 x_{1t} + \beta_2 x_{2t} + \beta_3 x_{3t} + \epsilon_t$$

- and test the joint hypothesis
 - $H_0 : \beta_2 = 4, \beta_3 = -2$
 - $H_a : \beta_2 \neq 4, \beta_3 \neq -2$
- One idea: variance-weighted squared distance of $(\hat{\beta}_2, \hat{\beta}_3)$ from $(4, -2)$
- Another idea: compare the variance of ϵ_t when $(\hat{\beta}_2, \hat{\beta}_3)$ are forced to be $(4, -2)$ vs. when they are freely estimated.
- These turn out to be the same thing: an F-test.

F-test review

Some board work.

Forecasting with one equation

- Suppose we want to forecast y_{1t} based on
 - $(y_{1,t-1}, \dots, y_{1,t-p})$
 - $(y_{2,t-1}, \dots, y_{2,t-p})$
 - ...
 - $(y_{n,t-1}, \dots, y_{n,t-p})$
 - and maybe some deterministic functions of the trend (time, time-squared, a sinusoid in time, seasonal dummies, etc.)
- Let $\mathbf{y}_t = (y_{1t}, y_{2t}, \dots, y_{nt})'$ be an $(n \times 1)$ vector
 - $\mathbf{x}_t = (1, \mathbf{y}'_{t-1}, \dots, \mathbf{y}'_{t-p})'$ is a $(k \times 1)$ vector, $k=np+1$
 - $\mathbf{y}_t, \mathbf{x}_t$ are covariance stationary

Forecasting with one equation

- Our forecast is $\hat{y}_{1t|t-1} = \beta'_{OLS} \mathbf{x}_t$
- For example: natural gas prices, oil prices, and drilling activity. Trying to forecast gas price returns.
 - y_{1t} = natural gas price return
 - y_{2t} = oil prices return
 - y_{3t} = oil & gas drilling activity

$$\begin{aligned}\hat{y}_{1t|t-1} = & \beta_0 + \beta_{1,ng}y_{1,t-1} + \dots + \beta_{p,ng}y_{1,t-p} \\ & + \beta_{1,oil}y_{2,t-1} + \dots + \beta_{p,oil}y_{2,t-p} + \beta_{1,dr}y_{3,t-1} + \dots + \beta_{p,dr}y_{3,t-p}\end{aligned}$$

Forecasting with one equation

- How to determine lag length
 - AIC as before
 - “Test down the model”: Successive F-tests of the p -th lag
 - H_0 : coefficients on $y_{1,t-p}, \dots, y_{n,t-p}$ are all jointly zero
- How to deal with MA terms
 - We will ignore them for now
 - We could model them explicitly
 - Tend to be captured by lags of y_2, y_3 , etc.

Forecasting with one equation

R Example: lag length in Henry Hub, WTI, drilling equation.

Forecasting with one equation

- Granger causality
 - H_0 : coefficients on $y_{2,t-1}, \dots, y_{2,t-p}$ are all zero (Granger causality of y_2)
 - We would say that oil price changes “Granger cause” natural gas price changes
 - This does not necessarily mean that oil price changes **cause** natural gas price changes. It means they have *forecasting information* about gas price changes.
 - In some cases, Granger causality can be the opposite of true causality.
 - ▶ Stock returns Granger-cause (predict) GDP growth, but *are caused by it*.
 - Likely no variable would Granger-cause oil prices, for example, but oil prices Granger-cause (and actually cause) many things.

Forecasting with one equation

R Example: Granger causality of WTI price and drilling activity on Henry Hub price.

Forecasting with one equation

- Structural stability
- Chow test:
 1. Estimate model over entire sample, save residual sum of squares RSS_0
 2. Estimate model with different coefficients before and after a date of suspected break t_1 :
 - 2.1 $y_{1t} = \mathbf{x}_t' \beta_1 (\delta_{[t \leq t_1]}) + \mathbf{x}_t' \beta_2 (\delta_{[t > t_1]}) + \epsilon_t$, save RSS_1
 3. Calculate $F(t_1) = \frac{(T-2k)(RSS_0 - RSS_1)}{kRSS_1}$, compare to $F(k, T - 2k)$
critical value, null is that $\beta_1 = \beta_2$ (stability)

Forecasting with one equation

- Andrews 1993 test:
 - Do Chow test for all t_1 between first 15% and last 15% of the sample.
 - Compare **largest** $kF(t_1)$ to critical value in Table 1 of Andrews (1993)
 - ▶ (nonstandard “Andrews Distribution” when comparing multiple dates - the largest draw from a sequence of Chi-squareds. Like an extreme value distribution).
 - Bai & Perron (1998) extend to multiple break points, select optimal number and timing of breaks.
 - ▶ R package “strucchange”.

Forecasting with one equation

R Example: structural break in the Brent oil price and LNG Asia price relationship

Vector Autoregression

- Suppose we are interested in more than just y_{1t} , or in potential feedbacks from y_{1t} through y_{2t} and y_{3t} .
- We have been estimating $y_{1t} = \pi_1' \mathbf{x}_t + \epsilon_{1t}$
 - $\mathbf{x}_t = (1, \mathbf{y}_{t-1}', \dots, \mathbf{y}_{t-p}')'$ is a $(k \times 1)$ vector, $k=np+1$
 - Here ϵ_{1t} is the forecast error for variable y_{1t} .
- We could run the analogous regression for $y_{2t} = \pi_2' \mathbf{x}_t + \epsilon_{2t}$ and so on.

Vector Autoregression

- Stack them in a vector system:

- $$\begin{bmatrix} y_{1t} \\ \vdots \\ y_{nt} \end{bmatrix} = \begin{bmatrix} \pi'_1 \\ \vdots \\ \pi'_n \end{bmatrix} \mathbf{x}_t + \begin{bmatrix} \epsilon_{1t} \\ \vdots \\ \epsilon_{nt} \end{bmatrix}$$

- $\mathbf{y}_t = \Pi' \mathbf{x}_t + \epsilon_t$ where Π' is $(n \times k)$

- $$\Pi' \mathbf{x}_t = \begin{bmatrix} c & \Phi_1 & \dots & \Phi_p \end{bmatrix} \begin{bmatrix} 1 \\ y_{t-1} \\ \vdots \\ y_{t-p} \end{bmatrix}$$

- $\mathbf{y}_t = \mathbf{c} + \Phi_1 \mathbf{y}_{t-1} + \dots + \Phi_p \mathbf{y}_{t-p} + \epsilon_t$ is called a vector autoregression (VAR).

Vector Autoregression

- Example: Henry Hub, WTI, drilling

$$\bullet \quad \begin{bmatrix} y_{ng,t} \\ y_{oil,t} \\ y_{dr,t} \end{bmatrix} = \begin{bmatrix} \beta_0^{ng} & \beta_{1,ng}^{ng} & \cdots & \beta_{3,dr}^{ng} \\ \beta_0^{oil} & \beta_{1,ng}^{oil} & \cdots & \beta_{3,dr}^{oil} \\ \beta_0^{dr} & \beta_{1,ng}^{dr} & \cdots & \beta_{3,dr}^{dr} \end{bmatrix} \begin{bmatrix} 1 \\ y_{ng,t-1} \\ \vdots \\ y_{dr,t-3} \end{bmatrix} + \begin{bmatrix} \epsilon_{ng,t} \\ \epsilon_{oil,t} \\ \epsilon_{dr,t} \end{bmatrix}$$

$$\bullet \quad \begin{bmatrix} y_{ng,t} \\ y_{oil,t} \\ y_{dr,t} \end{bmatrix} = \begin{bmatrix} \beta_0^{ng} \\ \beta_0^{oil} \\ \beta_0^{dr} \end{bmatrix} + \begin{bmatrix} \beta_{1,ng}^{ng} & \beta_{1,oil}^{ng} & \beta_{1,dr}^{ng} \\ \beta_{1,ng}^{oil} & \beta_{1,oil}^{oil} & \beta_{1,dr}^{oil} \\ \beta_{1,ng}^{dr} & \beta_{1,oil}^{dr} & \beta_{1,dr}^{dr} \end{bmatrix} \begin{bmatrix} y_{ng,t-1} \\ y_{oil,t-1} \\ y_{dr,t-1} \end{bmatrix} \\ + \dots + \begin{bmatrix} \beta_{3,ng}^{ng} & \beta_{3,oil}^{ng} & \beta_{3,dr}^{ng} \\ \beta_{3,ng}^{oil} & \beta_{3,oil}^{oil} & \beta_{3,dr}^{oil} \\ \beta_{3,ng}^{dr} & \beta_{3,oil}^{dr} & \beta_{3,dr}^{dr} \end{bmatrix} \begin{bmatrix} y_{ng,t-3} \\ y_{oil,t-3} \\ y_{dr,t-3} \end{bmatrix} + \begin{bmatrix} \epsilon_{ng,t} \\ \epsilon_{oil,t} \\ \epsilon_{dr,t} \end{bmatrix}$$

Vector Autoregression

- We can estimate this with conditional MLE like we learned before (conditional on the first p observations).
- Assume $\epsilon_t \sim N(\mathbf{0}, \Omega)$
- Treat the sample size as if its $T + p$ observations
- $\log L = -\frac{Tn}{2} \log(2\pi) - \frac{T}{2} \log|\Omega| - \frac{1}{2} \sum (\mathbf{y}_t - \Pi' \mathbf{x}_t)' \Omega^{-1} (\mathbf{y}_t - \Pi' \mathbf{x}_t)$
- Some really nice things happen when you maximize this:
- The i^{th} row of the MLE of Π' is the OLS estimate of the coefficients of the i^{th} equation.
 - We could estimate OLS equation by equation and get the exact same answers.

Vector Autoregression

- The residuals from the i^{th} row/equation are the OLS residuals from that equation, even if the system is estimated by MLE.
 - The MLE estimate of the variance Ω is equivalent to those OLS residuals in a straightforward way:
 - $\hat{\Omega} = \frac{1}{T} \sum \hat{\epsilon}_t \hat{\epsilon}_t'$

Vector Autoregression

- However testing joint hypotheses across equations is much easier using likelihood ratio test with Ω than F-tests.
 - The maximized value of the log likelihood is
$$-\frac{Tn}{2}(1 + \log 2\pi) - \frac{T}{2} \log |\hat{\Omega}|$$
 - Now if we want to test lag length for every equation in the whole system (p versus p-1 lags of each variable in each equation, in both cases using the last T of p+T observations because we burned the first p observations already), we have two estimates of Ω for two different lag lengths:
 - $\hat{\Omega}(p-1)$ and $\hat{\Omega}(p)$ that can be constructed from the OLS residuals from the two regressions with p and p-1 lags.
 - The likelihood ratio test statistic (twice the log likelihood ratio) is:
 - $T \left[\log |\hat{\Omega}(p-1)| - \log |\hat{\Omega}(p)| \right] \sim \chi^2(n^2)$
 - the small sample correction multiplies this by $T - k$, not T .

Vector Autoregression

- Other criteria for lag length selection include
 - Akaike information criterion (AIC): choose the specification with smallest AIC
 - ▶ $\log|\hat{\Omega}(p)| + 2p\frac{n^2}{T}$ penalizes the number of lags and equations per observation, rewards sample size
 - Schwarz or Schwarz-Bayes Criterion
 - ▶ $\log|\hat{\Omega}(p)| + \left(\frac{pn^2}{T}\right) \log T$, larger penalty on larger numbers of parameters relative to sample size.
 - Rules of thumb for macro data:
 - ▶ $p \geq 4$ for quarterly data
 - ▶ use lags 1-6 and 11-13 for monthly data

Impulse Response Functions, Forecasts

- We can rewrite the whole VAR $\mathbf{y}_t = \mathbf{c} + \Phi_1 \mathbf{y}_{t-1} + \dots + \Phi_p \mathbf{y}_{t-p} + \epsilon_t$ in companion form:

$$\bullet \begin{bmatrix} \mathbf{y}_t \\ \mathbf{y}_{t-1} \\ \vdots \\ \mathbf{y}_{t-p+1} \end{bmatrix}_{(np \times 1)} = \begin{bmatrix} \mathbf{c} \\ \mathbf{0} \\ \vdots \\ \mathbf{0} \end{bmatrix} + \begin{bmatrix} \Phi_1 & \Phi_2 & \dots & \Phi_p \\ \mathbf{I}_n & \mathbf{0} & \dots & \mathbf{0} \\ \vdots & \ddots & & \vdots \\ \mathbf{0} & \dots & \mathbf{I}_n & \mathbf{0} \end{bmatrix}_{(np \times np)} \begin{bmatrix} \mathbf{y}_{t-1} \\ \mathbf{y}_{t-2} \\ \vdots \\ \mathbf{y}_{t-p} \end{bmatrix} + \begin{bmatrix} \epsilon_t \\ \mathbf{0} \\ \vdots \\ \mathbf{0} \end{bmatrix}$$

Impulse Response Functions, Forecasts

- $\xi_t = \gamma + \mathbf{F}\xi_{t-1} + \mathbf{v}_t$
- $\xi_{t+2} = \gamma + \mathbf{F}\xi_{t+1} + \mathbf{v}_{t+2}$
- $\xi_{t+2} = (\mathbf{I}_{np} + \mathbf{F})\gamma + \mathbf{F}^2\xi_t + \mathbf{v}_{t+2} + \mathbf{Fv}_{t+1}$
- $$\xi_{t+s} = (\mathbf{I}_{np} + \mathbf{F} + \mathbf{F}^2 + \cdots + \mathbf{F}^{s-1})\gamma + \mathbf{F}^s\xi_t + \mathbf{v}_{t+s} + \mathbf{Fv}_{t+s-1} + \cdots + \mathbf{F}^{s-1}\mathbf{v}_{t+1}$$
- $\frac{\partial \xi_{t+s}}{\partial \xi_t'} = \mathbf{F}^s$

Impulse Response Functions, Forecasts

- We are really interested in the first n rows of ξ_{t+s} : $\frac{\partial \mathbf{y}_{t+s}}{\partial \mathbf{y}'_t}$
- $\mathbf{y}_{t+s} = \mathbf{c}_s + \epsilon_{t+s} + \Psi_1 \epsilon_{t+s-1} + \cdots + \Psi_{s-1} \epsilon_{t+1} + \Psi_s \mathbf{y}_t + \mathbf{F}_{12}^{(s)} \mathbf{y}_{t-1} + \cdots + \mathbf{F}_{1p}^{(s)} \mathbf{y}_{t-p+1}$
- $\Psi_s = \mathbf{F}_{11}^{(s)} = \frac{\partial \mathbf{y}_{t+s}}{\partial \mathbf{y}'_t} = \frac{\partial \mathbf{y}_{t+s}}{\partial \epsilon'_t}$
- Plotting Ψ_s as a function of s is called the “non-orthogonalized impulse response function.”

Impulse Response Functions, Forecasts

- Why “non-orthogonalized”? Errors across equations in the system are correlated.
 - $Cov(\epsilon_{it}\epsilon_{jt}) \neq 0$ in general.
 - Therefore a shock to ϵ_{it} contains information about all the other ϵ_{jt} 's, and so is not orthogonal to them in reality.
 - When we calculate the simple derivative, however, we hold the other ϵ_{jt} 's constant even though they are not.
 - We will learn another method to account for the additional information shortly.
- The row is the outcome being affected, and the column is the shock that's changing

Impulse Response Functions & Forecasts

- Row i , column j of Ψ_s gives $\frac{\partial y_{i,t+s}}{\partial y_{jt}} = \frac{\partial y_{i,t+s}}{\partial \epsilon_{jt}}$ holding constant ϵ_{kt} ($k \neq j$) and ϵ_{t+m} .
 - (even though ϵ_{kt} ($k \neq j$) not likely constant because of correlated shocks across equations).
- Notice that if our VAR(p) is stationary, it can be written as a VMA(∞):
 - $\mathbf{y}_{t+s} = \mu + \epsilon_{t+s} + \Psi_1 \epsilon_{t+s-1} + \Psi_2 \epsilon_{t+s-2} + \dots$
 - The Ψ 's have to be absolutely summable, so $\Psi_s \rightarrow \mathbf{0}$
 - Plots of these should have a decay
 - If not stationary, impulse response functions will have permanent effects.

Orthogonalized Impulse Response Functions

- If we are using impulse response functions to forecast, then a shock of ϵ_{1t} has information about contemporaneous shocks $\epsilon_{2t}, \dots, \epsilon_{nt}$
- We might want to take account of this extra information if we learn ϵ_{1t} but have not yet learned the value of the other shocks.
- Another way to ask the question:
 - Suppose we have only period $t - 1$ information $(\mathbf{y}_{t-1}, \mathbf{y}_{t-2}, \dots)$.
 - Then suddenly we learn the value of one variable, say y_{1t} .
 - How does learning only y_{1t} change our forecast of \mathbf{y}_{t+s} ?

Orthogonalized Impulse Response Functions

- Our VAR says:
- $\mathbf{y}_{t+s} = \mathbf{c}_s + \epsilon_{t+s} + \Psi_1 \epsilon_{t+s-1} + \cdots + \Psi_{s-1} \epsilon_{t+1} + \Psi_s \mathbf{y}_t + \mathbf{F}_{12}^{(s)} \mathbf{y}_{t-1} + \cdots + \mathbf{F}_{1p}^{(s)} \mathbf{y}_{t-p+1}$
- For forecasting, $E(\epsilon_{t+k}) = \mathbf{0}$, linear projection of \mathbf{y}_{t+s} on period t information is
- $\hat{\mathbf{y}}_{t+s|t} = \mathbf{c}_s + \Psi_s \mathbf{y}_t + \mathbf{F}_{12}^{(s)} \mathbf{y}_{t-1} + \cdots + \mathbf{F}_{1p}^{(s)} \mathbf{y}_{t-p+1}$
- Row i of this is our forecast of $y_{i,t+s}$ based on period t information.

Orthogonalized Impulse Response Functions

- What happens if we only knew period $t - 1$ information and then learned ϵ_{1t} ,
 - ϵ_{1t} is correlated with $\epsilon_{2t}, \dots, \epsilon_{nt}$ through the first column of Ω .
- Suppose our linear projection of \mathbf{y}_{t+s} is based only on knowing ϵ_{1t} and all the $t - 1$ information:
- $\hat{E}(\mathbf{y}_{t+s} | \epsilon_{1t}, \mathbf{y}_{t-1}, \mathbf{y}_{t-2}, \dots) =$
 $\mathbf{c}_s + \Psi_s \hat{E}(\epsilon_t | \epsilon_{1t}, \mathbf{y}_{t-1}, \mathbf{y}_{t-2}, \dots) + \mathbf{F}_{12}^{(s)} \mathbf{y}_{t-1} + \dots + \mathbf{F}_{1p}^{(s)} \mathbf{y}_{t-p+1}$
 - We use ϵ_{1t} to guess the rest of the period t information.

Orthogonalized Impulse Response Functions

- Let $\mathbf{a}_1 \epsilon_{1t} = \hat{E}(\epsilon_t | \epsilon_{1t})$
- $\hat{E}(\mathbf{y}_{t+s} | \epsilon_{1t}, \mathbf{y}_{t-1}, \mathbf{y}_{t-2}, \dots) = \mathbf{c}_s + \Psi_s \mathbf{a}_1 \epsilon_{1t} + \mathbf{F}_{12}^{(s)} \mathbf{y}_{t-1} + \dots + \mathbf{F}_{1p}^{(s)} \mathbf{y}_{t-p+1}$
- So the effect of y_{1t} on our forecast of \mathbf{y}_{t+s} if y_{1t} is all we know at period t is

$$- \frac{\partial \hat{E}(\mathbf{y}_{t+s} | \epsilon_{1t}, \mathbf{y}_{t-1}, \mathbf{y}_{t-2}, \dots)}{\partial y_{1t}} = \Psi_s \mathbf{a}_1$$

$$- \mathbf{a}_1 = \frac{\partial \hat{E}(\epsilon_t | \epsilon_{1t})}{\partial \epsilon_{1t}}$$

- How do we estimate this?

Orthogonalized Impulse Response Functions

- Notice that \mathbf{a}_1 accounts for how ϵ_{1t} is correlated with $\epsilon_{2t}, \dots, \epsilon_{nt}$ through the first column of Ω .
- Recall $\hat{\Omega} = \frac{1}{T} \sum \hat{\epsilon}_t \hat{\epsilon}_t'$
- We can factor this as $\hat{\Omega} = \hat{A} \hat{D} \hat{A}'$
- \hat{A} is a lower triangular matrix with ones on the principle diagonal, and \hat{D} is diagonal with all elements positive.
- $\hat{\mathbf{a}}_1$ is the first column of \hat{A}

Orthogonalized Impulse Response Functions

- Alternatively, we could take the Cholesky factorization of $\hat{\Omega}$:
 - $\hat{\Omega} = \hat{P}\hat{P}' = \hat{A}\hat{D}\hat{A}' = \hat{A}\hat{D}^{\frac{1}{2}}\hat{D}^{\frac{1}{2}}\hat{A}'$
 - \hat{P} is lower triangular with positive elements along principle diagonal.
 - Let $\hat{D}^{\frac{1}{2}}$ be the diagonal matrix whose diagonal is the diagonal of \hat{P} (and all other elements of \hat{D} are zero).
 - Then $\hat{A} = \hat{P}\hat{D}^{-\frac{1}{2}}$

Orthogonalized Impulse Response Functions

- The orthogonalized impulse response function is $\hat{\psi}_s \hat{\mathbf{a}}_1$.
 - Many researchers simply report $\hat{\psi}_s \hat{\mathbf{p}}_1$ where $\hat{\mathbf{p}}_1$ is the first column of $\hat{\mathbf{P}}$
 - This only differs by the scale.
 - $\hat{\psi}_s \hat{\mathbf{a}}_1$ is the effect of a one-unit increase in y_{1t}
 - $\hat{\psi}_s \hat{\mathbf{p}}_1$ is the effect of a one-standard deviation increase in y_{1t} .

Orthogonalized Impulse Response Functions

- Now what happens once I've learned y_{1t} , if I then learn y_{2t} but not y_{3t} , etc.?
 - $\hat{\psi}_s \hat{\mathbf{a}}_2$ where $\hat{\mathbf{a}}_2$ is the second column of $\hat{\mathbf{A}}$.
 - or use $\hat{\psi}_s \hat{\mathbf{p}}_2$ in terms of standard deviations.
- More generally, these are called “recursively orthogonalized impulse response functions”
 - $\hat{\mathbf{a}}_s \hat{\mathbf{A}}$
- The value of this is different depending on the ordering of the equations.
 - How you order them depends on what question you want to answer.
 - Typically order them in the order in which you think information would be made known.

Forecast Error Variance Decomposition

- How much of the variation in our forecast is coming from a particular variable?
- How much of the variance in our s -period ahead coal export forecast is coming from variance in the natural gas price?
- Define $\mathbf{u}_t = \mathbf{A}^{-1}\epsilon_t$.
- $E(\mathbf{u}_t\mathbf{u}_t') = \mathbf{A}^{-1}\Omega(\mathbf{A}^{-1})' = \mathbf{A}^{-1}\mathbf{A}\mathbf{D}\mathbf{A}'(\mathbf{A}^{-1})' = \mathbf{D}$
 - The variance covariance matrix of \mathbf{u}_t is diagonal.
 - \mathbf{u}_t are called “orthogonalized innovations”

Forecast Error Variance Decomposition

- Interpretation:

- $\epsilon_t = \mathbf{y}_t - \hat{\mathbf{y}}_{t|t-1}$
- $u_{1t} = \epsilon_{1t} = y_{1t} - \hat{E}(y_{1t} | \mathbf{y}_{t-1}, \mathbf{y}_{t-2}, \dots)$
- $u_{jt} = y_{jt} - \hat{E}(y_{jt} | y_{j-1,t}, y_{j-2,t}, \dots, y_{1t}, \mathbf{y}_{t-1}, \mathbf{y}_{t-2}, \dots)$

Forecast Error Variance Decomposition

- s-period-ahead forecast error:

- $\mathbf{y}_{t+s} - \hat{\mathbf{y}}_{t+s|t} = \epsilon_{t+s} + \Psi_1 \epsilon_{t+s-1} + \cdots + \Psi_{s-1} \epsilon_{t+1}$
- $= \mathbf{A} \mathbf{u}_{t+s} + \Psi_1 \mathbf{A} \mathbf{u}_{t+s-1} + \cdots + \Psi_{s-1} \mathbf{A} \mathbf{u}_{t+1}$
- We can break this down into the contribution from each u_{jt} :
- $\mathbf{A} \mathbf{u}_{t+m} = \mathbf{a}_1 u_{1,t+m} + \mathbf{a}_2 u_{2,t+m} + \cdots + \mathbf{a}_n u_{n,t+m}$

Forecast Error Variance Decomposition

- s-period-ahead mean squared error:

$$\begin{aligned}
 & - E(\mathbf{y}_{t+s} - \hat{\mathbf{y}}_{t+s|t})(\mathbf{y}_{t+s} - \hat{\mathbf{y}}_{t+s|t})' = \\
 & - = \Omega + \Psi_1 \Omega \Psi_1' + \Psi_2 \Omega \Psi_2' + \cdots + \Psi_{s-1} \Omega \Psi_{s-1}' \\
 & - = \mathbf{ADA} + \Psi_1 \mathbf{ADA} \Psi_1' + \Psi_2 \mathbf{ADA} \Psi_2' + \cdots + \Psi_{s-1} \mathbf{ADA} \Psi_{s-1}' \\
 & - = \sum_{m=0}^{s-1} \Psi_m (\mathbf{a}_1 \mathbf{a}_1' d_1 + \cdots + \mathbf{a}_n \mathbf{a}_n' d_n) \Psi_m' \\
 & - \text{The } k\text{'th term in paranthesis represents how much of our mean squared forecast error comes from not knowing}
 \end{aligned}$$

$u_{k,t+1}, u_{k,t+2}, \dots, u_{k,t+s}$:

- $$\begin{aligned}
 & - \sum_{m=0}^{s-1} \Psi_m (\mathbf{a}_k \mathbf{a}_k' d_k) \Psi_m' \\
 & - \text{You can plot this as a function of } s, \text{ similar to an impulse response function.} \\
 & - \text{Sometimes called “fevd” (which is the Stata subcommand for it) for “forecast error variance decomposition”.}
 \end{aligned}$$

Structural VAR

- VARs we have looked at so far have been “reduced form”: each equation in the system is a function of lags only: of its own value, and lags of other variables
- Maybe we can use/test structural information about contemporaneous relationships.
- This increases the number of parameters we need to estimate - only okay if the number of parameters is not too large to be recovered from reduced form estimates.
 - If the number of parameters is too large, we may impose some structural parameters in order to estimate others.

Structural VAR

- Consider a macroeconomic relationship
 - y_{1t} = real GDP growth
 - y_{2t} = inflation
 - y_{3t} = fed funds rate
 - y_{4t} = rate of growth of M2

Structural VAR

- Current spending (GDP growth) depends only on past shocks

$$- y_{1t} = k_1 + B_1^{(1,\cdot)} \mathbf{y}_{t-1} + \dots + u_{1t}$$

- Inflation has a Phillips curve relation between spending and inflation:

$$- y_{2t} = k_2 + b_0^{(2,1)} y_{1t} + B_1^{(2,\cdot)} \mathbf{y}_{t-1} + \dots + u_{2t}$$

- Federal reserve responds to current output growth and inflation:

$$- y_{3t} = k_3 + b_0^{(3,1)} y_{1t} + b_0^{(3,2)} y_{2t} + B_1^{(3,\cdot)} \mathbf{y}_{t-1} + \dots + u_{3t}$$

- Money demand depends on current output, inflation, and interest rates:

$$- y_{4t} = k_4 + b_0^{(4,1)} y_{1t} + b_0^{(4,2)} y_{2t} + b_0^{(4,3)} y_{3t} + B_1^{(4,\cdot)} \mathbf{y}_{t-1} + \dots + u_{4t}$$

Structural VAR

- We can stack these equations into a vector dynamic structural model:
 - $\mathbf{B}_0 \mathbf{y}_t = \mathbf{k} + \mathbf{B}_1 \mathbf{y}_{t-1} + \cdots + \mathbf{B}_p \mathbf{y}_{t-p} + \mathbf{u}_t$
 - Any dynamic structural model that has a linear approximation can be written in this way.
 - $$\mathbf{B}_0 = \begin{bmatrix} 1 & 0 & 0 & 0 \\ b_0^{(2,1)} & 1 & 0 & 0 \\ b_0^{(3,1)} & b_0^{(3,2)} & 1 & 0 \\ b_0^{(4,1)} & b_0^{(4,2)} & b_0^{(4,3)} & 1 \end{bmatrix}$$
 - The lower triangular form is a special case for this model

Structural VAR

- If we premultiply by \mathbf{B}_0^{-1} , we get the reduced form:

- $\mathbf{y}_t = \mathbf{c} + \Phi_1 \mathbf{y}_{t-1} + \dots + \Phi_p \mathbf{y}_{t-p} + \epsilon_t$
 $n \times 1$
- $\mathbf{c} = \mathbf{B}_0^{-1} \mathbf{k}$
- $\Phi_j = \mathbf{B}_0^{-1} \mathbf{B}_j$
- $\epsilon_t = \mathbf{B}_0^{-1} \mathbf{u}_t$

Structural VAR

- The reduced form VAR is just another representation of a linearized dynamic structural model.
 - If we knew the structural parameters \mathbf{B} and shocks \mathbf{u} we could calculate the VAR parameters Φ and ϵ and the associated impulse response functions.
 - $\frac{\partial \mathbf{y}_{t+m}}{\partial \mathbf{u}'_t} = \frac{\partial \mathbf{y}_{t+m}}{\partial \epsilon'_t} \frac{\partial \epsilon_t}{\partial \mathbf{u}'_t} = \Psi_m \mathbf{B}_0^{-1}$

Structural VAR

- So can we estimate the reduced form and recover the structural parameters?
 - Reduced form VAR parameters: $n + n(n + 1)/2 + n^2 p$
 - Structural VAR parameters: $n + n(n + 1)/2 + n^2 p + n^2$
 - Need enough restrictions on \mathbf{B}_0 (e.g., zeros for some elements) and $E(\mathbf{u}_t \mathbf{u}_t')$ (e.g., diagonal matrix) in order for their to be a one-to-one mapping.
 - The earlier lower triangular form, with diagonal $E(\mathbf{u}_t \mathbf{u}_t')$, is one such example.

Structural VAR

- If there is a one-to-one mapping, then estimate reduced form and figure out the appropriate transformation of Φ and Ω
 - For example in the case above:
 - $\epsilon_t = \mathbf{B}_0^{-1} \mathbf{u}_t$
 - $\Omega = \mathbf{B}_0^{-1} \mathbf{D} (\mathbf{B}_0^{-1})'$
 - Find triangular factorization of $\hat{\Omega}$ as before.
 - When the structural equation is lower triangular, this is the same as the recursively orthogonalized impulse response function.

Structural VAR

- For non-recursive structural models:
 - Let $E(\mathbf{u}_t \mathbf{u}_t') = \mathbf{D}$
 - Estimate reduced form by OLS, put enough restrictions on $\mathbf{D}, \mathbf{B}_0, \dots, \mathbf{B}_p$ so that there is a unique mapping from $\Omega, \Phi_1, \dots, \Phi_p$.
 - Typical approach: let $\mathbf{B}_1, \dots, \mathbf{B}_p$ be unrestricted.
 - $E(\epsilon_t \epsilon_t') = \Omega = \mathbf{B}_0^{-1} \mathbf{D} (\mathbf{B}_0^{-1})'$
 - $\hat{\Omega}$ has $n(n+1)/2$ distinct elements. so we can have just that many unknowns in \mathbf{B}_0 and \mathbf{D}
 - Need to impose as many structural restrictions from theory, long run equilibrium, etc., to reduce parameters to a number that can be recovered from the reduced form (or estimated directly using MLE).

Scalar Case

- $y_t = \phi y_{t-1} + w_t$
- $y_{t+1} = \phi y_t + w_{t+1} = \phi^2 y_{t-1} + \phi w_t + w_{t+1}$
- $y_{t+m} = \phi^{m+1} y_{t-1} + \sum_{s=0}^m \phi^s w_{t+m-s}$
- If shock w_t is larger, but w_{t+j} is unchanged, what happens to y_{t+m} at date $t + m$?
- $\frac{\partial y_{t+m}}{\partial w_t} = \phi^m$

Scalar Case

- $\frac{\partial y_{t+m}}{\partial w_t} = \phi^m$
- Plot this derivative as a function of m , should be decaying in m if $\phi < 1$.
 - This is an “impulse response function”.
 - If $\phi < 0$ we have oscillation.
 - If $|\phi| > 1$, we have explosive unstable effect.
 - If $\phi = 1$, we have a permanent effect.

Vector Case

- \mathbf{y}_t = vector of n variables observed at t (system of n equations).
- $\mathbf{y}_t = \mathbf{F}\mathbf{y}_{t-1} + \mathbf{w}_t$, \mathbf{F} is an $n \times n$ matrix of coefficients.
- $\mathbf{y}_{t+1} = \mathbf{F}\mathbf{y}_t + \mathbf{w}_{t+1} = \mathbf{F}^2\mathbf{y}_{t-1} + \mathbf{F}\mathbf{w}_t + \mathbf{w}_{t+1}$, where $\mathbf{F}^2 = \mathbf{F} \cdot \mathbf{F}$
- If $\mathbf{F} = \begin{bmatrix} \phi_1 & \phi_2 \\ 1 & 0 \end{bmatrix}$, then
$$\mathbf{F}^2 = \begin{bmatrix} \phi_1 & \phi_2 \\ 1 & 0 \end{bmatrix} \begin{bmatrix} \phi_1 & \phi_2 \\ 1 & 0 \end{bmatrix} = \begin{bmatrix} \phi_1^2 + \phi_2 & \phi_1\phi_2 \\ \phi_1 & \phi_2 \end{bmatrix}$$
- $\mathbf{y}_{t+m} = \mathbf{F}^{m+1}\mathbf{y}_{t-1} + \mathbf{F}^m\mathbf{w}_t + \mathbf{F}^{m-1}\mathbf{w}_{t+1} + \dots + \mathbf{F}\mathbf{w}_{t+m-1} + \mathbf{w}_{t+m}$
$$= \mathbf{F}^{m+1}\mathbf{y}_{t-1} + \sum_{s=0}^m \mathbf{F}^s\mathbf{w}_{t+m-s}$$

Vector Case

- What happens if j^{th} component of \mathbf{w}_t (w_{jt}) increases by 1 unit, with all other w fixed? What would happen to i^{th} component of \mathbf{y}_{t+m} ($y_{i,t+m}$)?
- $\frac{\partial y_{i,t+m}}{\partial w_{jt}} =$ the row i , column j element of the $n \times n$ matrix $\frac{\partial \mathbf{y}_{t+m}}{\partial \mathbf{w}'_t} = \mathbf{F}^m$.
- For example, to find the effect of w_{1t} on $y_{1,t+3}$, find the (1,1) element of

$$- \mathbf{F}^3 = \begin{bmatrix} \phi_1^2 + \phi_2 & \phi_1 \phi_2 \\ \phi_1 & \phi_2 \end{bmatrix} \begin{bmatrix} \phi_1 & \phi_2 \\ 1 & 0 \end{bmatrix}$$

$$- \implies F_{1,1}^3 = \phi_1^3 + 2\phi_1\phi_2$$

Vector Case

- We might want to know how \mathbf{F}^m behaves as a function of m generally.

- Depends on the eigenvalues of \mathbf{F} , or values of λ that make

$$|\mathbf{F} - \lambda \mathbf{I}_n| = 0$$

- or
$$\begin{vmatrix} f_{11} - \lambda & \cdots & f_{1n} \\ \vdots & \ddots & \vdots \\ f_{n1} & \cdots & f_{nn} - \lambda \end{vmatrix} = 0$$

Vector Case

- Calculating this determinant using cofactor expansion gives an n^{th} -order polynomial in λ
 - $a_0 + a_1\lambda + a_2\lambda^2 + \dots + a_n\lambda^n = 0$
 - Suppose this polynomial has n distinct roots. Then there exists a nonsingular $n \times n$ matrix \mathbf{T} such that we can diagonalize

► $\mathbf{F} = \mathbf{T}\mathbf{\Lambda}\mathbf{T}^{-1}$, where $\mathbf{\Lambda} = \begin{bmatrix} \lambda_1 & 0 & \dots & 0 \\ 0 & \lambda_2 & & \\ \vdots & & \ddots & \\ 0 & & & \lambda_n \end{bmatrix}$

- If there are not n distinct roots, we can use the Jordan normal form of \mathbf{F} and get similar results.

Vector Case

- Now $\mathbf{F}^2 = (\mathbf{T}\mathbf{\Lambda}\mathbf{T}^{-1})(\mathbf{T}\mathbf{\Lambda}\mathbf{T}^{-1}) = \mathbf{T}\mathbf{\Lambda}^2\mathbf{T}^{-1}$ or $\mathbf{F}^m = \mathbf{T}\mathbf{\Lambda}^m\mathbf{T}^{-1}$, where

$$- \mathbf{\Lambda}^m = \begin{bmatrix} \lambda_1^m & \cdots & 0 \\ \vdots & \ddots & \\ 0 & & \lambda_n^m \end{bmatrix}$$

- We want to say something about the nature of the impulse response coefficients $\frac{\partial y_{i,t+m}}{\partial w_{jt}}$ as m grows.
 - ▶ This is now the i, j element of $\mathbf{T}\mathbf{\Lambda}^m\mathbf{T}^{-1}$, which because of \mathbf{T} has the form $a_{ij}^{(1)}\lambda_1^m + \dots + a_{ij}^{(n)}\lambda_n^m$
 - ▶ Knowing the eigenvalues of \mathbf{F} and the elements of \mathbf{T} and \mathbf{T}^{-1} tells us the impulse response.

Vector Case

- If the λ' s are all real:
 - the impulse response function is a weighted average of the possible impulse response functions for the first-order scalar case.
 - If $|\lambda_s| < 1 \forall s$, then effects will die out over time (system is stable)
 - ▶ $\lim_{m \rightarrow \infty} \frac{\partial y_{i,t+m}}{\partial w_{jt}} = 0 \forall i, j$
 - If $|\lambda_s| > 1$ for some s , then the system is explosive.
 - If $\lambda_n = 1$ and $|\lambda_s| < 1$ for $s = 1, \dots, n-1$, then $\frac{\partial y_{i,t+m}}{\partial w_{jt}} \rightarrow a_{ij}^{(n)}$ and there is a permanent effect (even if w_n is not the shock).

Vector Case

- If the λ 's are complex:

- For example, $\mathbf{F} = \begin{bmatrix} \phi_1 & \phi_2 \\ 1 & 0 \end{bmatrix}$

- $\Rightarrow |\mathbf{F} - \lambda \mathbf{I}| = \begin{vmatrix} \phi_1 - \lambda & \phi_2 \\ 1 & -\lambda \end{vmatrix} = \lambda^2 - \phi_1 \lambda - \phi_2 = 0$

- $\lambda = \frac{\phi_1 \pm \sqrt{\phi_1^2 + 4\phi_2}}{2}$, so the roots can be complex conjugates if $\phi_1^2 + 4\phi_2 < 0$

- $\lambda_1 = a + bi$, $\lambda_2 = a - bi$ for some a, b .

Vector Case

- If the λ' s are complex:
 - $\frac{\partial \mathbf{y}_{t+m}}{\partial \mathbf{w}_t'} = \mathbf{F}^m$ is still real, but the matrices in the decomposition $\mathbf{T}\mathbf{\Lambda}^m\mathbf{T}^{-1}$ has complex elements.
 - The $\sqrt{-1}$ will multiply and cancel out.
 - ▶ Complex numbers are only part of the intermediate analysis to get some insight.
 - ▶ Want to ask how does λ^m behave as a function of m for $\lambda = a + bi$

Vector Case, complex roots

- Recall that we can write $a + bi = R(\cos\theta + i \cdot \sin\theta) = Re^{i\theta}$, where the last equality comes from Euler's formula.
 - $R = \sqrt{a^2 + b^2}$
 - ▶ $\lambda^m = (a + bi)^m = R^m e^{i\theta m} = R^m (\cos(\theta m) + i \cdot \sin(\theta m))$
 - ▶ λ^m behaves like R^m multiplied by a sinusoidal function.
 - Magnitude of effect depends on size or $R = \sqrt{a^2 + b^2}$
 - ▶ $R < 1$ means effects die out over time (as m increases)
 - ▶ $R > 1$ means effects explode
 - ▶ $R = 1$ means there are permanent effects, and permanent oscillations in this case.

Summary

- stability of $\mathbf{y}_t = \mathbf{F}\mathbf{y}_{t-1} + \mathbf{w}_t$ depends on eigenvalues of \mathbf{F} .
 - If $\|\lambda_s\| < 1$ for all $s = 1, \dots, n$ (or the modulus - absolute value when real, R when complex)
 - ▶ $|\lambda_s| < 1$ when real, $R_s = \sqrt{a^2 + b^2} < 1$ when complex, or if all λ_s inside the unit circle
 - ▶ Then the system is **stable**.

2nd-order scalar difference equation

- $y_t = \phi_1 y_{t-1} + \phi_2 y_{t-2} + w_t$
- Define $\mathbf{y}_t = \begin{bmatrix} y_t \\ y_{t-1} \end{bmatrix}$
- $\begin{bmatrix} y_t \\ y_{t-1} \end{bmatrix} = \begin{bmatrix} \phi_1 & \phi_2 \\ 1 & 0 \end{bmatrix} \begin{bmatrix} y_{t-1} \\ y_{t-2} \end{bmatrix} + \begin{bmatrix} w_t \\ 0 \end{bmatrix}$
- $\mathbf{y}_t = \mathbf{F} \mathbf{y}_{t-1} + \mathbf{w}_t$
- Call this the “companion form”.
- We have already analyzed this as a first order vector system.

p^{th} -order scalar difference equation

- $y_t = \phi_1 y_{t-1} + \phi_2 y_{t-2} + \dots + \phi_p y_{t-p} + w_t$
- Companion form:

$$\bullet \mathbf{y}_t = \begin{bmatrix} y_t \\ y_{t-1} \\ \vdots \\ y_{t-p+1} \end{bmatrix}, \mathbf{F} = \begin{bmatrix} \phi_1 & \phi_2 & \cdots & \phi_{p-1} & \phi_p \\ 1 & 0 & & & 0 \\ 0 & 1 & & & 0 \\ \vdots & & \ddots & & \\ 0 & 0 & \cdots & 1 & 0 \end{bmatrix}, \mathbf{w}_t = \begin{bmatrix} w_t \\ 0 \\ \vdots \\ 0 \end{bmatrix}$$

- $\mathbf{y}_t = \mathbf{F}\mathbf{y}_{t-1} + \mathbf{w}_t$
- $\frac{\partial y_{1,t+m}}{\partial w_{1t}} = c_1 \lambda_1^m + \dots + c_p \lambda_p^m$ if all λ 's distinct.

p^{th} -order scalar difference equation

- Eigenvalues of \mathbf{F} satisfy

$$\lambda^p - \phi_1 \lambda^{p-1} - \phi_2 \lambda^{p-2} - \dots - \phi_p = 0$$

- effects die out if $\|\lambda_s\| < 1$ for all $s = 1, \dots, p$
- effects explode if $\|\lambda_s\| > 1$ for some s
- effects are permanent if $\|\lambda_s\| = 1$ for some s .

Spurious Regression

- Suppose we have two independent random walks
 - $\Delta y_{1t} = \epsilon_{1t}, \Delta y_{2t} = \epsilon_{2t}$
 - $\epsilon_t \sim iid(\mathbf{0}, \Omega)$ where $\Omega = \begin{bmatrix} \sigma_1^2 & 0 \\ 0 & \sigma_2^2 \end{bmatrix}$

Spurious Regression

- Suppose we regress one random walk on the other
 - $y_{1t} = \alpha + \gamma y_{2t} + u_t$
 - true $\alpha = \gamma = 0$
 - However, $\hat{\gamma}_{ols}$ and $\hat{\alpha}_{ols}$ have nonstandard limiting distributions and are not mean zero.
 - Also, the t-stats do not have a limiting distribution, so there are no correct critical values.
 - Also, $R^2 = 1 - \frac{SSE}{SST}$ which is a ratio of variances of random walks. This converges to a nonstandard distribution which doesn't tell us much about the true variation explained.
- spurious regression: estimated relationship looks awesome when there is really nothing.

Spurious Regression

- How to recognize and characterize a spurious regression:
 - The error in a spurious regression will **not** be $I(0)$ stationary.
 - $y_t = x_t' \beta + u_t$
 - If I can find a β that makes $u_t \sim I(0)$, then this is **not** spurious.
 - If no such β exists, then it is spurious.
 - Spurious case:
 - $y_{1t} = \alpha + \gamma y_{2t} + u_t$
 - true $\alpha = \gamma = 0 \implies u_t = y_{1t} = \epsilon_{11} + \epsilon_{12} + \dots + \epsilon_{1t}$
 - $u_t = y_{1t} \sim I(1)$.

Spurious Regression

- How to avoid a spurious regression
 1. Difference all variables before regression, or otherwise make sure all series are stationary.
 - ▶ $\Delta y_{1t} = \alpha + \gamma \Delta y_{2t} + u_t$
 - ▶ $\hat{\alpha}, \hat{\gamma}$ in this case converge normally.
 2. Include lags of both LHS and RHS variables as additional regressors (ADL model).
 - ▶ $y_{1t} = \alpha + \gamma y_{2t} + \phi y_{1,t-1} + \delta y_{2,t-1} + u_t$
 - ▶ true $\alpha = \gamma = \delta = 0, \phi = 1, u_t = \epsilon_t = \Delta y_{1t}$

Spurious Regression

- How to avoid a spurious regression
 1. Difference all variables before regression, or otherwise make sure all series are stationary.
 2. Include lags of both LHS and RHS variables as additional regressors (ADL model).
 - ▶ $y_{1t} = \alpha + \gamma y_{2t} + \phi y_{1,t-1} + \delta y_{2,t-1} + u_t$
 - ▶ To analyze this we need to **rotate** the regressors
 - ▶ $y_{1t} = \alpha + \gamma \Delta y_{2t} + \phi y_{1,t-1} + \delta^* y_{2,t-1} + u_t$
 - ▶ $\delta^* = \delta + \gamma$
 - ▶ Δy_{2t} is mean-zero, stationary, while $y_{1,t-1}$ and $y_{2,t-1}$ are not.

Spurious Regression

- How to avoid a spurious regression
 1. Difference all variables before regression, or otherwise make sure all series are stationary.
 2. Include lags of both LHS and RHS variables as additional regressors (ADL model).
 - ▶ In this setup all estimates are consistent (we're not in spurious case)
 - ▶ Now γ converges to a normal, can use standard t-test.
 - ▶ Other parameters converge to something non-normal, can still use t-stat tested against alternative distributions.

Spurious Regression

- How to avoid a spurious regression
 1. Difference all variables before regression, or otherwise make sure all series are stationary.
 2. Include lags of both LHS and RHS variables as additional regressors (ADL model).
 - ▶ In original regression, γ and δ are now normal, can use standard t-stats.
 - ▶ Intuition for δ : $\delta = \delta^* - \gamma$. δ^* was consistent in rotated regression, converges to itself, while γ converges to a Normal.
 - ▶ F-tests of joint hypotheses are **not** standard.
 - ▶ t-test of $\phi = 1$ is nonstandard in same way as Dickey-Fuller test.
- Conclusion: Including lags can help fix nonstationarity, but if you know they're nonstationary, just difference them first.

Unit Roots in Vector Autoregressions

- true model: $\Delta \mathbf{y}_t = \zeta_1 \Delta \mathbf{y}_{t-1} + \dots + \zeta_{p-1} \Delta \mathbf{y}_{t-p+1} + \epsilon_t$
 - the data **should** have been differenced. (Case 2)
 - Stationarity assumption for this model in VAR looks like
 - $|\mathbf{I}_n - \zeta_1 z - \zeta_2 z^2 - \dots - \zeta_{p-1} z^{p-1}| = 0$ with $\|z\| < 1 \implies \Delta \mathbf{y}_t \sim I(0)$
 - ▶ VAR(p-1) with no drift.

Unit Roots in Vector Autoregressions

- Suppose we did not difference, estimated model VAR in levels with a constant.
- $\mathbf{y}_t = \alpha + \Phi_1 \mathbf{y}_{t-1} + \Phi_2 \mathbf{y}_{t-2} + \dots + \Phi_p \mathbf{y}_{t-p} + \epsilon_t$
- As in previous Dickey-Fuller case, this is a rotation of
- $\mathbf{y}_t = \alpha + \rho \mathbf{y}_{t-1} + \zeta_1 \Delta \mathbf{y}_{t-1} + \dots + \zeta_{p-1} \Delta \mathbf{y}_{t-p+1} + \epsilon_t$
- True $\alpha = \mathbf{0}$ and $\rho = \mathbf{I}_n$
- As before, $\Phi_p = -\zeta_{p-1}$,
- $\Phi_s = \zeta_s - \zeta_{s-1}$ for $s = 2, \dots, p-1$
- $\Phi_1 = \rho + \zeta_1$

Unit Roots in Vector Autoregressions

- As before, all the ζ_s 's are asymptotically normal, can use standard t-tests and F-tests.
- $\hat{\alpha}$ and $\hat{\rho}$ will be nonstandard but consistent (converge to something).
- Implications for estimating the levels regression:
 - t-tests and F-tests of elements of Φ_s are asymptotically valid
 - confidence intervals of impulse-response functions are asymptotically valid.
 - What's not okay: testing joint hypotheses that involve ρ .
 - For example, in testing for Granger causality, we might want to test that a particular element in \mathbf{y}_{t-s} for $s = 1, \dots, p$ has predictive power, or joint test of $\Phi_1^{(2,1)} = \Phi_2^{(2,1)} = \dots = \Phi_p^{(2,1)} = 0$.
 - ▶ But this involves ρ and ζ , so we can't do it.
- On the other hand, if we should NOT difference, and we do, our model is misspecified - you will wash out important relationships, including cointegration.

Cointegration

- Let's illustrate with a special case:
- Suppose
 - $y_{1t} \sim I(1)$
 - $\Delta y_{2t} = u_{2t}$ with u_{2t} white noise (so y_{2t} is difference stationary).
 - $y_{1t} = \gamma y_{2t} + u_{1t}$
 - $\mathbf{u}_t = \begin{bmatrix} u_{1t} \\ u_{2t} \end{bmatrix}$, $E(\mathbf{u}_t) = 0$, $E(\mathbf{u}_t \mathbf{u}_t') = \begin{bmatrix} \sigma_1^2 & 0 \\ 0 & \sigma_2^2 \end{bmatrix}$ if $t = s$

Cointegration

- Then $\Delta y_{1t} = \gamma \Delta y_{2t} + \Delta u_{1t} = \gamma u_{2t} + u_{1t} - u_{1,t-1} \sim I(0)$
 - This is a first-order moving average, with u_{2t} stationary.
- In other words, $y_{1t}, y_{2t} \sim I(1)$ but $y_{1t} - \gamma y_{2t} \sim I(0)$.
 - $[1 \quad -\gamma]'$ is the **cointegrating vector**.

Cointegration

Definition: an $(n \times 1)$ vector \mathbf{y}_t is said to be cointegrated if each element y_{it} is $I(1)$ but there is a linear combination $\alpha' \mathbf{y}_t \sim I(0)$, where in general α is called the cointegrating vector.

- This is not spurious regression, even though we're regressing random walk on a random walk.
 - In spurious case, there is **no** choice of α that can make the vector stationary.
 - Spurious: $u_{1t} \sim I(1)$
 - Cointegrated: $u_{1t} \sim I(0)$

Cointegration

- In our example, y_{1t} will inherit the random walk that y_{2t} follows, but the cointegrating relationship keeps them close together.
 - They deviate by u_{1t} , which if $I(0)$ will always return to a fixed mean.
 - This is an appealing model of long run market (or ecological?) relationships.

Cointegration

- Rewrite this as a vector system:

$$- \begin{bmatrix} 1 & -\gamma \\ 0 & 1 \end{bmatrix} \begin{bmatrix} y_{1t} \\ y_{2t} \end{bmatrix} = \begin{bmatrix} 0 & 0 \\ 0 & 1 \end{bmatrix} \begin{bmatrix} y_{1t-1} \\ y_{2t-1} \end{bmatrix} + \begin{bmatrix} u_{1t} \\ u_{2t} \end{bmatrix}$$

$$- \text{premultiply by } \begin{bmatrix} 1 & -\gamma \\ 0 & 1 \end{bmatrix}^{-1} = \begin{bmatrix} 1 & \gamma \\ 0 & 1 \end{bmatrix}$$

$$- \begin{bmatrix} y_{1t} \\ y_{2t} \end{bmatrix} = \begin{bmatrix} 0 & \gamma \\ 0 & 1 \end{bmatrix} \begin{bmatrix} y_{1t-1} \\ y_{2t-1} \end{bmatrix} + \begin{bmatrix} u_{1t} + \gamma u_{2t} \\ u_{2t} \end{bmatrix}$$

- Let $\begin{bmatrix} \epsilon_{1t} \\ \epsilon_{2t} \end{bmatrix} = \begin{bmatrix} u_{1t} + \gamma u_{2t} \\ u_{2t} \end{bmatrix}$, which is white noise because the u 's are white noise.

Cointegration

- $E(\epsilon_t) = \mathbf{0}$ and $E(\mathbf{e}_t \mathbf{e}_t') = \begin{bmatrix} \sigma_1^2 + \gamma^2 \sigma_2^2 & \gamma \sigma_2^2 \\ \gamma \sigma_2^2 & \sigma_2^2 \end{bmatrix}$ if $t = s$, $\mathbf{0}$ otherwise.
- So we could run this vector system in levels.
- $\mathbf{y}_t = \Phi \mathbf{y}_{t-1} + \mathbf{e}_t$, where $\Phi = \begin{bmatrix} 0 & \gamma \\ 0 & 1 \end{bmatrix}$

Cointegration

- Notice a few things

1. the vector \mathbf{y}_t has a unit root (is nonstationary).

$$1.1 \quad |\mathbf{I}_2 - \Phi z| = \begin{vmatrix} 1 & -\gamma z \\ 0 & 1 - z \end{vmatrix} = 1 - z = 0 \text{ when } z = 1.$$

2. we could write the system in differences as (analogous to our rotated regression).

- ▶ $\begin{bmatrix} \Delta y_{1t} \\ \Delta y_{2t} \end{bmatrix} = \begin{bmatrix} -1 & \gamma \\ 0 & 0 \end{bmatrix} \begin{bmatrix} y_{1t-1} \\ y_{2t-1} \end{bmatrix} + \begin{bmatrix} u_{1t} + \gamma u_{2t} \\ u_{2t} \end{bmatrix}$
- ▶ $\Delta \mathbf{y}_t = \rho \mathbf{y}_{t-1} + \mathbf{e}_t$
- ▶ this is a hybrid of changes and levels.

Cointegration

- $\Delta \mathbf{y}_t = \rho \mathbf{y}_{t-1} + \mathbf{e}_t$
1. If we should have had a system (VAR) in changes, then $\underline{\rho}$ will be $\mathbf{0}$.
 2. If we should have had a system (VAR) in levels (y_{1t}, y_{2t} were stationary), then $\underline{\rho}$ is an arbitrary set of coefficients and has full rank.
 3. If the system is cointegrated, $\underline{\rho}$ has rank 1 and can be written as an outer product of two vectors, one of which is the cointegrating vector.

$$- \underline{\rho} = \begin{bmatrix} -1 \\ 0 \end{bmatrix} \begin{bmatrix} 1 & -\gamma \end{bmatrix} = \mathbf{b}\alpha'$$

Cointegration

- Why would we want to write the system this way: $\Delta \mathbf{y}_t = \rho \mathbf{y}_{t-1} + \mathbf{e}_t$?
- Notice we have $\Delta \mathbf{y}_t = \mathbf{b} \alpha' \mathbf{y}_{t-1} + \mathbf{e}_t$, and $\alpha' \mathbf{y}_{t-1} = u_{1,t-1}$
 - So $\Delta \mathbf{y}_t = \mathbf{b} u_{1,t-1} + \mathbf{e}_t$,
 - Today's changes $\Delta \mathbf{y}_t$ depend on how far y_1 and y_2 deviated from their cointegrating relation last period.
 - Because y_1 and y_2 are random walks, $\Delta \mathbf{y}_t$ should be white noise.
 - Indeed it is, but with a special kind of white noise: $u_{1,t-1}$

Cointegration

- In general: An $(n \times 1)$ vector \mathbf{y}_t is said to exhibit $h < n$ cointegrating relations if
 1. each element of $\mathbf{y}_t \sim I(1)$
 2. there exists an $(n \times h)$ matrix \mathbf{A} of full rank h such that each element of the $(h \times 1)$ vector $\mathbf{A}'\mathbf{y}_t \sim I(0)$
 - The rows of \mathbf{A}' are called the cointegrating relations or vectors.
- In other words, if $n > 2$, there might be more than one linear combination which is cointegrated.

Cointegration

- Implications for general case:
 - If we can write \mathbf{y}_t as a vector system (VAR) then that system will have a unit root:
 1. $\mathbf{y}_t = \alpha + \Phi_1 \mathbf{y}_{t-1} + \dots + \Phi_p \mathbf{y}_{t-p} + \mathbf{e}_t$
 2. $|\mathbf{I}_n - \Phi_1 z - \dots - \Phi_p z^p| = 0$ when $z = 1$

Cointegration

- Implications for general case:
 - If previous slide holds, then we can write it as a rotated regression:
 1. $\Delta \mathbf{y}_t = \alpha + \rho \mathbf{y}_{t-1} + \zeta_1 \Delta \mathbf{y}_{t-1} + \dots + \zeta_{p-1} \Delta \mathbf{y}_{t-p+1} + \mathbf{e}_t$
 2. $\underline{\rho} = \mathbf{B}\mathbf{A}'$ where \mathbf{B} is an $n \times h$ matrix of rank h , \mathbf{A}' is an $h \times n$ matrix whose rows are the cointegrating vectors.
 3. $\underline{\rho}$ has rank $h < n$. As before, $\underline{\rho} = \mathbf{0}$ implies the VAR should be in differences $\Delta \mathbf{y}_t$ and if $\underline{\rho}$ has rank n , the VAR should be in levels \mathbf{y}_t
 4. the $\Delta \mathbf{y}$'s are stationary, and the \mathbf{y} 's are nonstationary. The rows of \mathbf{A}' take exactly the linear combinations of the levels to not ruin the stationarity of the remaining differences.

Cointegration

- Implications for general case:
 - If previous slide holds, then we can think of $\mathbf{A}'\mathbf{y}_t = \mathbf{z}_t$ as the stationary residual from the cointegrating relations of \mathbf{y}_t which is often called the “error correction term” and write
 1. $\Delta\mathbf{y}_t = \alpha + \mathbf{B}\mathbf{z}_{t-1} + \zeta_1\Delta\mathbf{y}_{t-1} + \dots + \zeta_{p-1}\Delta\mathbf{y}_{t-p+1} + \mathbf{e}_t$
 2. Notice that this is just a VAR in changes with an extra term (that would have been in the error if we didn't include it).
 3. \mathbf{z}_{t-1} was the degree of deviation from the cointegrating relation last period. \mathbf{B} measures how quickly we return to it.
- Conclusion: with a cointegrated system, we can estimate it as a VAR in levels, or as a VAR in differences with the error correction term.

Estimation and testing of a single cointegrating relation

- Let $\underline{\alpha}' = [1 \quad -\alpha_2 \quad \dots \quad -\alpha_n]$ be the cointegrating vector.
- $\underline{\alpha}' \mathbf{y}_t \sim I(0)$ for the true α but $\sim I(1)$ for any other α .
- Consider OLS of y_{1t} on y_{2t}, \dots, y_{nt}
- $y_{1t} = \alpha_2 y_{2t} + \dots + \alpha_n y_{nt} + u_t$
- $\min_{\underline{\alpha}} \sum (y_{1t} - \alpha_2 y_{2t} - \dots - \alpha_n y_{nt})^2$

Estimation and testing of a single cointegrating relation

- Recall from spurious regression that $\frac{1}{T} \sum u_t^2 \rightarrow V$ if u_t stationary, and $\rightarrow \infty$ if not.
 - If T is large enough, we will be able to tell the difference from something going to infinity.
- Conclusion: OLS of y_{1t} on remaining y 's is a good way to estimate $\underline{\alpha}$ and find out if it's really cointegrated or not.
- If it's not cointegrated, this will be a spurious regression.
- To tell the difference, take
 - H_0 : no cointegration (spurious regression, residuals from OLS regression have a unit root).
 - Notice this is the same null hypothesis as the DF test for stationarity, which we can apply to the residuals (however, testing distribution is different).
 - If they are stationary, reject the null and conclude the series are cointegrated.

Estimation and testing of a single cointegrating relation

- Procedure (for example):

1. Estimate by OLS

1.1 $y_{1t} = \alpha + \gamma_2 y_{2t} + \dots + \gamma_n y_{nt} + u_t$

1.1.1 includes a constant to account for any possible drift.

1.2 save the residuals \hat{u}_t

2. Estimate by OLS

2.1 $\hat{u}_t = \rho \hat{u}_{t-1} + \zeta_1 \Delta \hat{u}_{t-1} + \dots + \zeta_{p-1} \Delta \hat{u}_{t-p+1} + \nu_t$

2.2 no constant (Case 1).

2.3 If spurious, $\rho = 1$, fail to reject DF test null

2.4 If cointegrated, $\rho < 1$, reject DF test null.

2.5 **caveat:** because the null involves a spurious regression, the distribution of ρ will have a **different nonstandard distribution** here than it does in the DF test.

2.6 This is called a Phillips-Ouliaris-Hansen test - same setup as Dickey-Fuller, including different cases for drift and trend, but with different distributions.

Estimation and testing of a single cointegrating relation

- Estimating and testing of more than one cointegrating relation can be done using an MLE-type procedure called Johansen's algorithm,
 - this can be executed using the “vec” command in Stata with option , rank(#) for number of cointegrating vectors.
 - To figure out number of cointegrating vectors, use statistics derived from the Johansen MLE algorithm
 - ▶ Johansen's maximum eigenvalue, and Johansen's trace method tell you the number of cointegrating vectors.
 - ▶ Or use Stata's “vecrank” command.